Scaling of matrix-product-state plus cosine-filter ground-state preparation: gap, overlap and target in the J_1–J_2 model, with a comparison to QETU

Authors to be added
(Dated: November 2, 2026)

A common route to a ground state on a quantum computer compiles a classical trial state and removes its residual excited weight with a quantum filter. We benchmark this route on the spin-$\frac{1}{2}$ J_1–J_2 model, with a snake matrix-product state (MPS) as the trial and the one-ancilla cosine filter of Stetcu, Baroni and Carlson (SBC), by exact simulation of the Trotterized circuits on up to 16 qubits (20 for one test) and with planned circuits beyond. Three findings stand out. (i) The a-priori certificate's overhead is almost all in its Trotter bound: the planned filter is within 2.0× of the length simulation needs, but the bounded step is 60–75 times finer than needed, and with a step transferred from the next-smaller system the planned filter costs 0.13–0.18 of its certified cost and met the target in all 112 tested cases and 30 of 30 at 20 qubits. (ii) Filtering pays only where the trial is hard to improve. On chains it wins only as a last-mile polisher, and those wins disappear when the MPS alone may use any bond dimension. On the 4×4 cylinder, filtering a cheap snake MPS beats compiling a better snake MPS or exact generic state preparation at 13 of 14 targets measured, even against a lower bound on the MPS circuits; the implementable version, with the transferred step, wins 5 of 14 with constructive MPS costs but 0 against that lower bound with any bond dimension, and none certified; on a 4×5 cylinder (20 qubits) it wins 11 and 2 of 14. The measured advantage is robust; the implementable one is fragile, though it grows from 4×4 to 4×5. (iii) Against a Chebyshev (QETU) filter with its longest admissible query, SBC wins at $\epsilon \approx 10^{-2}$, ties at 10^{-3} and loses from 10^{-4}. A controlled study that holds the gap (0.30–0.55, a large-gap regime) and the overlap fixed separately at $N = 8$–14 shows that the gap sets the circuit depth and the overlap the repetitions. The infidelity bound held in all 96 simulated circuits, all at $N \leq 14$ with exact inputs; for larger systems it is conditional on DMRG estimates of the gap and overlap, and a 10–25% gap overestimate breaks the optimized and Chebyshev plans; error bars taken from the DMRG bond-dimension ladder covered the exact gap in all 75 tests.

I. INTRODUCTION

A common route to a ground state on a quantum computer starts classically: an approximate method produces a trial state that can be compiled into a short circuit, and a quantum subroutine removes its residual excited-state weight. For one-dimensional systems the natural classical engine is DMRG [3, 4], which yields the ground-state energy E_0, the gap and a matrix-product state (MPS) whose low-bond-dimension truncations compile exactly into staircase circuits [5]. The quantum step can be phase estimation [6], a polynomial filter implemented by quantum singular-value transformation [7–9], or a product of cosine filters implemented with a single ancilla [1, 13]. The last option, the projection algorithm of Stetcu, Baroni and Carlson (SBC) [1], needs little beyond time evolution: each pulse is one controlled time evolution between ancilla rotations and a measurement, the filter is designed classically, and the success probability is known in advance.

This paper asks three questions about the SBC filter on a concrete, physically relevant model, the open J_1–J_2 Heisenberg chain:

1. *What does it cost?* Two-qubit gates, arbitrary rotations, evolution time and repetitions, both as certified by an a-priori bound and as measured by exact simulation of the Trotterized circuit.

2. *What sets the cost?* The spectral gap Δ and the trial overlap γ normally change together with system size N and with the coupling. We decouple them: the gap is held fixed across N by tuning J_2, and the overlap is held fixed by a one-parameter rotation of the trial circuit, so each can be varied while the other stays put.

3. *When does it pay off?* The alternative to filtering a cheap trial state is compiling a better one. We compare the two routes at equal target infidelity ϵ.

The answers rest on a modular pipeline (Sec. II) that takes energy estimates and an ansatz for any Pauli-sum Hamiltonian, designs a cosine filter with a certified suppression ratio, chooses Trotter step counts from commutator bounds, and returns an a-priori bound on the infidelity together with a gate count. Section III introduces the model and validates the bound. Section IV is the core of the paper: the controlled gap/overlap study, which yields the cost laws and explains them through a factorization into filter length, Trotter step and gates per step. Section V follows the natural chains to $N = 48$, tests robustness to wrong inputs and measures where the filter beats a compiled MPS. Section VI repeats that comparison on $4 \times L_x$ cylinders, where it comes out the other way, and Sec. VII compares with a Chebyshev polynomial filter (QETU) [9] under identical accounting. Section VIII sets every measured comparison beside its certified counterpart, and Sec. IX tests two rules that choose the Trotter step without the ground state, including one size beyond every other simulation. The comparison of the two SBC schedules, cost maps, and an analysis of the

Chebyshev filter's degree are in Appendices B and C.

Relation to prior work. The filter is SBC's [1]; it shares its circuit with the rodeo algorithm [13, 20], which draws the times at random. Our closed-form "dyadic" schedule, a halving ladder of times, is the one SBC and Rule *et al.* use for quantum-number projection [1, 21], and for energy it is a finite form of the rodeo super-iteration of Cohen and Oh [22], recently generalized to other geometric ratios [23]; what we add is its certified bound on the scaled spectral window and its measured cost. The optimized schedule is SBC's energy recipe with a certified constraint. The Trotter bounds are those of Childs *et al.* [2], and the certificate follows the standard argument for polynomial filters [7, 9]. QETU has been costed with Trotterized evolution for lattice gauge theories [24] and run on small Hubbard models [25, 26]. Resource studies that compile an MPS or other classical trial and then refine it exist for quantum chemistry [11, 12, 27], and Watanabe *et al.* [28] stop an MPS encoder early and finish with probabilistic imaginary-time evolution on Heisenberg chains. For the two-dimensional J_1–J_2 model, the quantum–classical crossover has been estimated for phase estimation against DMRG with state preparation taken as negligible [29]. What is new here is (i) a controlled separation of gap and overlap at fixed size, (ii) a measured, not only bounded, account of every resource and its factorization, which shows where the certificate's overhead sits, (iii) a like-for-like measured comparison of cosine and Chebyshev filters, including the Chebyshev filter's query length, (iv) a map of where filtering a cheap MPS beats compiling a better one or preparing the state generically, in one and two dimensions, and (v) a test of a-posteriori Trotter-step rules that turns part of the measured advantage into an implementable one. The method is not new; the contribution is the benchmark and what it shows about where the cost of this workflow sits. The code is the `sbpipe` package, available with all data at `https://github.com/aaronseymour7/FilterPipeline`; every number below is produced by the scripts named in Appendix D.

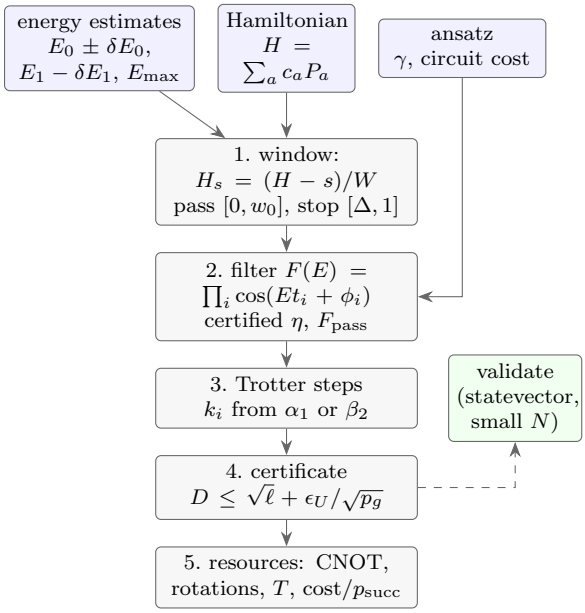


FIG. 1. The pipeline. Blue: inputs, which any classical method can supply. Grey: the stages of `refine()`. Green: the optional statevector check used in Sec. III B.

II. THE SBC FILTER AND THE RESOURCE PIPELINE

A. Inputs and the spectral window

The Hamiltonian is $H = \sum_a c_a P_a$ with real c_a and Pauli strings P_a, in a fixed order that is also the Trotter order, and with an optional grouping of terms (here: one group per bond). The pipeline assumes

$$E_0 \in [E_0^{\text{est}} - \delta E_0,\ E_0^{\text{est}} + \delta E_0], \qquad E_1 \ge E_1^{\text{est}} - \delta E_1, \qquad E_{\text{top}} \le E_{\max}, \tag{1}$$

where E_1 is the lowest level above the ground space that the trial state can overlap with. When the trial state's quantum numbers are unknown, as for the DMRG trial states below, which are not symmetry-adapted, this is the global first excited level; restricting to the trial state's symmetry sector could only raise it. The certificate is a statement about the actual Hamiltonian only if (1) holds for it: with an overestimated E_1 a real excitation can sit in the filter's transition region, and the bound then says nothing (Sec. V B). If no $E_{\max}$ is given, the Weyl bound $E_{\max} = \sum_g \lambda_{\max}(h_g)$ over the term groups h_g is used; it is rigorous and, for the antiferromagnetic J_1–J_2 chain, attained by the ferromagnetic state. With $s = E_0^{\text{est}} - \delta E_0$ and $W = E_{\max} - s$, the scaled Hamiltonian $H_s = (H - s)/W$ has its ground energy in the *pass band* $[0, w_0]$, $w_0 = 2\delta E_0/W$, and every other relevant level in the *stop band* $[\Delta, 1]$, $\Delta = (E_1^{\text{est}} - \delta E_1 - s)/W$.

The ansatz enters through γ, a lower estimate of its weight on the ground space, and through the gate cost of the circuit that prepares it. A statevector is needed only for validation.

B. The filter and three schedules

One SBC pulse [1] applies $\mathsf{H}_a\, e^{-itH_s \otimes Z_a}\, R_z^{(a)}(2\phi)\, \mathsf{H}_a$ to the system and an ancilla in $|0\rangle$; the ancilla-$|0\rangle$ block of this unitary is $\cos(H_s t + \phi)$. (SBC write the same block with $e^{-i(tH+\delta)\otimes Y_a}$.) Keeping outcome 0 after m pulses applies

$$F(H_s) = \prod_{i=1}^{m} \cos(H_s t_i + \phi_i), \qquad T = \sum_i |t_i|. \tag{2}$$

Two numbers characterize a filter: the pass-band floor $F_{\text{pass}} \le \min_{[0,w_0]} |F|$ and the suppression ratio $\eta = \sup_{[\Delta,1]} |F|/F_{\text{pass}}$. Both are certified, not sampled: since $|F'| \le L := \sum_i |t_i|$ and $|F''| \le L^2$, an interval branch-and-bound gives a rigorous upper bound on $\sup |F|$ over the stop band, and $F_{\text{pass}} = |F(0)| - Lw_0$ (the certifier is `sbpipe.floor.certify`; it is used for the optimized schedule, the other two have closed-form certificates). Given a target leakage ϵ_ℓ and γ, the design target is $\eta^2 \le \gamma\epsilon_\ell/[(1-\epsilon_\ell)(1-\gamma)]$ (Lemma 1). We provide three schedules:

Uniform: m equal pulses $t = \pi/(1+\Delta)$, $\phi = 0$, so $F = \cos^m(Et)$. For $E \in [\Delta, 1]$ we have $Et \in [\theta, \pi-\theta]$, $\theta = \pi\Delta/(1+\Delta)$, hence $\sup_{[\Delta,1]} |\cos Et| = \cos\theta$. Then $m \approx 2\ln(1/\eta)/\theta^2$ and $T \approx 2\ln(1/\eta)/(\pi\Delta^2)$: quadratic in the inverse gap. It is the baseline that any design should beat.

Dyadic: r repetitions of the halving ladder $t = \tau, 2\tau, \ldots, 2^{K-1}\tau$, $\tau = \pi/(1+\Delta)$, $\phi = 0$. This is the schedule SBC use for quantum-number projection [1, 21], and for the energy it is a finite rodeo super-iteration [22] (whose infinite form gives $\sin x/x$). One block is a Dirichlet kernel,

$$\prod_{k=0}^{K-1} \cos(2^k E\tau) = \frac{\sin(2^K E\tau)}{2^K \sin(E\tau)}, \tag{3}$$

which equals 1 at $E = 0$. The identity is exact for every E (repeated use of $\sin 2x = 2\sin x \cos x$). The suppression bound is a separate step that uses the spectral window: for $E \in [\Delta, 1]$, $E\tau \in [\theta, \pi-\theta]$, so $\sin(E\tau) \ge \sin\theta$ and, with $|\sin(2^K E\tau)| \le 1$, the block is at most $\eta_b = 1/(2^K \sin\theta)$ in absolute value there. This bound relies on the scaling H_s, which places every relevant level in $[0, 1]$, and on the choice of τ; it does not hold outside $[\Delta, 1]$. With $r = \lceil \ln\eta / \ln\eta_b \rceil$ blocks, $T = r\tau(2^K - 1)$; choosing $2^K \sin\theta \approx e$ gives

$$T_{\text{dyadic}} \approx \frac{e\,\tau}{\sin\theta} \ln\frac{1}{\eta} \;\approx\; \frac{e}{\Delta} \ln\frac{1}{\eta}, \tag{4}$$

the optimal $O(\Delta^{-1}\log\eta^{-1})$ scaling [7] with an explicit constant and no optimization; Cohen and Oh reach the same conclusion for super-iterations on a generic gapped spectrum [22]. The pipeline searches K and r for the smallest T. Both closed-form schedules are certified by the inequalities just given (with $F_{\text{pass}} = \cos^m(w_0 t)$ and $1 - Lw_0$, respectively), which avoids interval arithmetic on filters with $L \sim 10^3$–10^5; the test suite checks them against a 2×10^6-point grid.

Optimized: times and phases chosen to maximize $F(0)^2$ subject to the certified constraint on η at fixed (T, m), over a grid of $x = T\Delta/\pi$ and $m \in \{4, 6, 8\}$, picking the pair with the smallest $T^2/F(0)^2$. This is SBC's recipe for energy filters [1] with a certified constraint in place of a grid objective; it uses the SLSQP designer `sbpipe/floor.py`. Nonzero phases trade success probability ($F(0) < 1$) for a shorter T.

For each schedule the leakage share $\epsilon_\ell/\epsilon \in \{0.5, 0.25, 0.1, 0.05\}$ is scanned and the plan with the smallest expected CNOT count per success is kept. An "auto" mode does the same over the dyadic and optimized schedules together.

C. Trotterization

Each controlled evolution $e^{-it_i H_s \otimes Z_a}$ is approximated by k_i steps of the first-order (Lie) or second-order (Strang) product formula over the non-identity terms; the identity term of H_s becomes a single ancilla R_z per pulse and is exact. For $B_a = \sum_{b>a} H_b$ the bounds of Childs *et al.* [2, Props. 15–16] are

$$\|S_1(t) - e^{-itH}\| \le \tfrac{t^2}{2}\alpha_1, \quad \alpha_1 = \textstyle\sum_a \|[B_a, H_a]\|, \tag{5}$$

$$\|S_2(t) - e^{-itH}\| \le t^3\beta_2, \quad \beta_2 = \tfrac{1}{12}\textstyle\sum_a \|[B_a,[B_a,H_a]]\| + \tfrac{1}{24}\textstyle\sum_a \|[H_a,[H_a,B_a]]\|, \tag{6}$$

so k steps over time t cost $\alpha_1 t^2/(2k)$ or $\beta_2 t^3/k^2$. The filter times t_i are in units of the scaled Hamiltonian, so the constants are evaluated for the terms of H_s, i.e. for coefficients c_a/W (the identity shift does not enter any commutator). Equivalently $\alpha_1(H_s) = \alpha_1(H)/W^2$ and $\beta_2(H_s) = \beta_2(H)/W^3$, with physical evolution time t_i/W; the code uses the first convention, and we checked that it agrees with the second to machine precision. Tensoring every term with Z_a leaves α_1 and β_2 unchanged because $Z_a^2 = 1$. For Pauli strings $[P_b, P_a] = 2P_bP_a$ when they anticommute and 0 otherwise, so each commutator is a short Pauli sum supported on a few qubits around term a; we evaluate its norm exactly on that support (dense for ≤ 8 qubits, Lanczos to a relative tolerance of 10^{-12} up to 12 with a 10^{-9} safety factor, triangle inequality beyond, which never occurred here). The Lanczos values are numerically exact but not certified upper bounds; replacing them by the triangle inequality would make every constant rigorous at a higher cost. The constants are maximized over the forward and reversed term order, so they hold for either convention. Every pulse evolves under the same $H_s \otimes Z_a$, so the per-pulse bounds $\alpha_1 t_i^2/(2k_i)$ and $\beta_2|t_i|^3/k_i^2$ share one constant; the phase ϕ_i enters only the ancilla rotation, and the bound is symmetric in the sign of t_i. With equal constants, the continuous relaxation of $\min \sum_i k_i$ subject to $\sum_i(\text{bound}_i) \le \epsilon_T$ is solved by $k_i \propto |t_i|$ for both orders (Lagrange). We take $k_i = \lceil |t_i|/\delta t \rceil$ with the common step δt of that solution. Rounding up keeps the bound, costs at most m steps more than the continuous optimum and need not be the integer optimum; the budget it frees goes to rotation synthesis. The certified times are untouched.

D. The certificate

Theorem 1. *Assume (1) and that the trial state has ground-space weight at least γ. Let the designed filter have certified η and F_{pass}, let ϵ_U bound the operator-norm error of the implementation (Trotter error $\sum_i \epsilon_i$ plus $R\delta$ for the R arbitrary rotations of the trial and filter circuits, each synthesized to accuracy δ), and let*

$$\ell = \frac{(1-\gamma)\eta^2}{\gamma + (1-\gamma)\eta^2}, \qquad p_g = \gamma F_{\mathrm{pass}}^2. \tag{7}$$

Let $|g\rangle = \Pi_0|\psi\rangle/\|\Pi_0|\psi\rangle\|$ be the normalized projection of the trial state onto the ground space, with Π_0 the ground-space projector; $|g\rangle$ is a pure state also when the ground space is degenerate. For a noiseless implementation the state $|\tilde{\psi}_{\mathrm{out}}\rangle$ prepared conditioned on the ancilla outcomes 0 is pure, and

$$D \le \sqrt{\ell} + \frac{\epsilon_U}{\sqrt{p_g}}, \qquad 1 - \mathrm{Tr}(\Pi_0\tilde{\rho}) \le 1 - |\langle g|\tilde{\psi}_{\mathrm{out}}\rangle|^2 = D^2, \tag{8}$$

where D is the trace distance between $\tilde{\rho} = |\tilde{\psi}_{\mathrm{out}}\rangle\langle\tilde{\psi}_{\mathrm{out}}|$ and $|g\rangle$. Moreover $\langle H\rangle - E_0 \le (E_{\mathrm{top}} - E_0)D^2$, and the success probability is at least $(\sqrt{p_g} - \epsilon_U)^2$. Below, $1-F$ denotes $1-|\langle g|\tilde{\psi}_{\mathrm{out}}\rangle|^2$; for the nondegenerate ground state of the chains studied here it is the usual infidelity.

The proof (Appendix A) combines a leakage lemma for the exact filter with the observation that a vector error ϵ_U on an unnormalized vector of norm at least $\sqrt{p_g}$ moves its direction by at most $\epsilon_U/\sqrt{p_g}$. The pipeline sets the budget $\epsilon_U \le (\sqrt{\epsilon} - \sqrt{\ell})\sqrt{p_g}$, gives 90% of it to the product formula and the rest to rotation synthesis. If the trial state already satisfies $1-\gamma \le \epsilon$, the pipeline also offers the empty filter ($\eta = 1$, no post-selection, $\ell = 1-\gamma$) and keeps whichever plan is cheaper.

E. Resource model

All-to-all connectivity and the Clifford+R_z gate set are assumed. Every Trotter factor of a pulse is controlled by the ancilla. For a Heisenberg bond (i, j) the three commuting terms form one exact three-qubit block,

$$e^{-ict(X_iX_j+Y_iY_j+Z_iZ_j)\otimes Z_a} = V^\dagger e^{-ict(Z_iZ_a+Z_jZ_a-Z_iZ_jZ_a)}V,$$
$$V = \mathsf{H}_i\,\mathrm{CNOT}_{ij}, \tag{9}$$

since the Bell-basis change V maps XX, YY, ZZ to $Z_i, -Z_iZ_j, Z_j$. The diagonal factor is a phase polynomial on the parities $i \oplus a$, $i \oplus j \oplus a$, $j \oplus a$, applied by $\mathrm{CNOT}_{ia}\,R_z\,\mathrm{CNOT}_{ja}\,R_z\,\mathrm{CNOT}_{ia}\,R_z\,\mathrm{CNOT}_{ja}$ on the ancilla wire [34]. One controlled bond thus costs 6 CNOTs and 3 R_z, against 12 CNOTs for three separate weight-3 Pauli rotations and 8 for qiskit's `PauliEvolutionGate` synthesis (both checked in the test suite, as is Eq. (9)); it is the controlled analogue of the 3-CNOT two-qubit gate [32, 33]. A numerical search over 300 four-CNOT layouts found no exact circuit. Since $X_iX_j + Y_iY_j + Z_iZ_j = 2\,\mathrm{SWAP}_{ij} - 1$, the block is a power of SWAP in ancilla-Z form, $e^{-2ict\,\mathrm{SWAP}_{ij}\otimes Z_a}$ up to an ancilla phase, and the comparison with known circuits is with controlled powers of SWAP. The ordinary controlled form (identity when the ancilla is 0), as a Hadamard-test or standard rodeo circuit would use it, puts the phase on the singlet indicator times the ancilla, a doubly controlled phase in the Bell basis, which needs all seven parities and by the same construction costs 6 CNOTs plus the 2 of V; for SWAP itself (the Fredkin gate) 7 CNOTs, or 5 general two-qubit gates [39], are known. The ancilla-Z form that the cosine filter needs uses only the three parities that contain the ancilla, which is why it needs 6. The block is the same unitary as its three term rotations, so the Trotterized circuit and its error are those of the term-wise product formula. The cheaper control by an anticommuting Pauli string, used by controlled-reversal rodeo circuits [10] and control-free QETU [9], is not available: no product of single-qubit Paulis anticommutes with all of XX, YY and ZZ on a bond. A Lie step costs one block per bond; k Strang steps cost $k(2B-2)+1$ blocks for B bonds, because the two middle half-steps and the outer half-steps of adjacent steps merge. Each pulse adds one ancilla R_z. An arbitrary rotation at accuracy δ costs $3\log_2(1/\delta)$ T gates, the leading term of Ross–Selinger synthesis [14]. The cost of one attempt is the trial circuit plus the filter, and the expected cost per success is the attempt cost divided by the certified lower bound $(\sqrt{p_g} - \epsilon_U)^2$ on the success probability of the implemented circuit from Theorem 1 (repeat-until-success; an early-abort protocol that measures after each pulse would be cheaper and is not modeled). The pipeline spends the whole budget $\epsilon_U = (\sqrt{\epsilon} - \sqrt{\ell})\sqrt{p_g}$, so this denominator is $p_g(1 - \sqrt{\epsilon} + \sqrt{\ell})^2 \ge (1 - \sqrt{\epsilon})^2 p_g$, i.e. at least $0.81\,p_g$ at $\epsilon = 10^{-2}$. Without a filter nothing is post-selected and every attempt succeeds. Alongside each T count we report the number of arbitrary rotations per attempt and their common synthesis accuracy δ.

III. MODEL, INPUTS AND VALIDATION

A. Model and front end

We study the open spin-$\frac{1}{2}$ chain

$$H = J_1 \sum_{i=1}^{N-1} \mathbf{S}_i \cdot \mathbf{S}_{i+1} + J_2 \sum_{i=1}^{N-2} \mathbf{S}_i \cdot \mathbf{S}_{i+2}, \tag{10}$$

$J_1 = 1$, at $J_2 = 0$ (Heisenberg), $J_2 = 0.2411$ (close to the critical point of the gapless-to-dimerized transition [16, 17]) and $J_2 = 0.4$ (dimerized), for $N = 8$ to 48. Its Pauli form has $B = N-1$ bonds ($3B$ terms of weight 2) at $J_2 = 0$ and $B = 2N-3$ otherwise. Terms are ordered by bond: nearest-neighbor bonds with even then odd left site, then next-nearest-neighbor bonds likewise.

The front end (`sbpipe/frontends/dmrg_quimb.py`) builds the MPO from the Pauli sum and runs two-site DMRG (quimb [18]) with bond dimensions up to 256 ($N \le 24$) or 128 ($N > 24$) and truncation cutoff 10^{-12}. E_1 is the lowest state of $H + \mu(S^z_{\rm tot} - 1)^2$, $\mu = 2$, i.e. the lowest triplet, which we use as the global first excitation. The error bars $\delta E_0, \delta E_1$ are the energy change between the last two bond dimensions, floored at 10^{-10}; they are convergence heuristics, not bounds. Both DMRG energies are variational upper bounds, so a residual DMRG error in E_1 (or convergence of the penalty run to a higher state) biases the gap upward, which is the dangerous direction (Sec. V B); these error bars do not protect against that bias, and the certificate for $N > 14$ is conditional on the gap. The trial state is the variationally optimal MPS at fixed bond dimension $\chi \in \{1, 2, 4, 8\}$, and γ is its overlap with the reference MPS. This is an estimate of the overlap with the exact ground state, not a lower bound on it, so for $N > 14$ the certificate is equally conditional on γ. An MPS of bond dimension 2^b is exactly a staircase of $N - b$ unitaries on $b + 1$ qubits; we cost each with the quantum Shannon decomposition bound $\frac{23}{48}4^n - \frac{3}{2}2^n + \frac{4}{3}$ CNOTs (3, 20, 100 for $n = 2, 3, 4$) [15] and $4^n - 1$ rotations.

For $N \le 14$ we compare with exact diagonalization (ED). DMRG reproduces E_0 to 2.3×10^{-11}, the first excited level of any symmetry to 2.4×10^{-11} (so the first excitation is the triplet in every ED case), and the trial overlaps to 4.0×10^{-8} (largest difference, for a $\chi = 1$ trial). The Weyl bound equals the exact top of the spectrum to 1.2×10^{-14}. Table I lists the inputs. The trial overlap decays roughly exponentially with N at fixed χ, because a fixed bond dimension captures a fixed amount of entanglement per bond while the truncation error accumulates over N bonds, so the filter has more to do at larger N.

TABLE I. DMRG inputs: ground energy, gap $E_1 - E_0$ and trial overlap γ of the optimal bond-dimension-χ MPS. Source: `results/inputs.json`.

N	J_2	E_0	gap	$\gamma_{\chi=1}$	γ_2	γ_4	γ_8
8	0	-3.374933	0.3927	0.134	0.775	0.999	1.0000
12	0	-5.142091	0.2809	0.060	0.591	0.994	1.0000
16	0	-6.911737	0.2193	0.028	0.475	0.984	0.9999
24	0	-10.453786	0.1530	0.006	0.328	0.954	0.9994
32	0	-13.997316	0.1177	0.001	0.234	0.913	0.9978
48	0	-21.085956	0.0809	0.000	0.125	0.817	0.9902
8	0.2411	-3.135001	0.3859	0.113	0.904	1.000	1.0000
12	0.2411	-4.738077	0.2720	0.046	0.806	1.000	1.0000
16	0.2411	-6.343316	0.2100	0.019	0.703	0.998	1.0000
24	0.2411	-9.556221	0.1443	0.003	0.517	0.989	0.9997
32	0.2411	-12.770432	0.1099	0.001	0.373	0.972	0.9989
48	0.2411	-19.200243	0.0744	0.000	0.191	0.920	0.9946
8	0.4	-3.025191	0.4037	0.036	0.976	1.000	1.0000
12	0.4	-4.545626	0.2956	0.003	0.945	1.000	1.0000
16	0.4	-6.066666	0.2371	0.000	0.911	0.999	1.0000
24	0.4	-9.109325	0.1748	0.000	0.839	0.995	0.9998
32	0.4	-12.152237	0.1418	0.000	0.768	0.990	0.9995
48	0.4	-18.238247	0.1075	0.000	0.640	0.976	0.9981

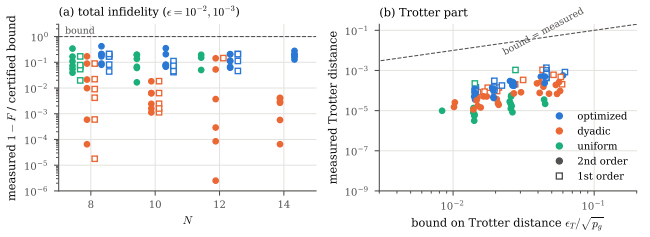


FIG. 2. All simulated plans. (a) Measured infidelity of the planned circuit divided by its certified bound; points below the dashed line respect the bound and are offset horizontally by schedule and order. (b) Measured distance between the Trotterized and the exact filtered state against its certified bound. Points below the dashed lines respect the bound. Filled circles: second order; open squares: first order.

B. Validation of the certificate ($N \le 14$)

For $N \in \{8, 10, 12, 14\}$, the three couplings, $\epsilon \in \{10^{-2}, 10^{-3}\}$, the three schedules and both Trotter orders we plan a circuit for the $\chi = 2$ trial state and, when its size permits, simulate it exactly: the ancilla-0 block of the Trotterized circuit is $\frac{1}{2}(e^{-i\phi}S(t) + e^{i\phi}S(-t))$ per pulse, applied to the trial statevector and compared with the ED ground state. Of 144 plans, 96 were simulated (the others, mostly uniform or first-order plans with 10^4–10^8 steps, exceed our simulation budget), and the certificate held in all 96 (Fig. 2).

The bound is conservative by a median factor 14 (range 2.4–4.0 $\times 10^5$) in infidelity. Splitting it shows where: the leakage term ℓ overestimates the exact filter's infidelity by a median 4.3× (range 1.1–1.3 $\times 10^5$). By schedule the medians are 1.4 (optimized), 4 (uniform) and 54 (dyadic): the optimized design pushes $|F|$ to $\eta F_{\rm pass}$ across the band, whereas the closed-form dyadic bound $1/(2^K \sin\theta)$ is reached only near the band edge. More generally, η is a worst case over the whole stop band while the trial state's excited weight sits mostly low in the band. The Trotter term overestimates the measured Trotter distance by a median 185× (range 26–4.5 $\times 10^3$): commutator bounds are worst-case over all states, while the filtered state lives in the low-energy subspace, where Trotter errors are much smaller [19, 31], and empirical Trotter errors are known to lie far below the bounds [30]. The measured success probability equals p_g to two digits in every case, because after filtering almost all of the surviving norm is ground-state weight.

Table II compares the schedules for one case. The uni-

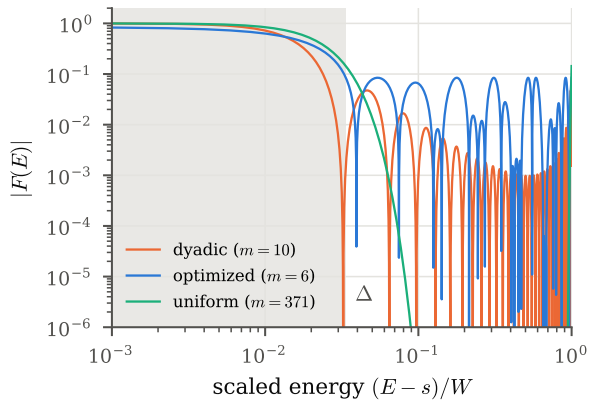


FIG. 3. $|F(E)|$ of the three schedules for $N = 12$, $J_2 = 0.2411$, $\epsilon = 10^{-2}$, second order. Dotted lines: certified ηF_{pass}. Grey: the gap region. The horizontal axis is linear below Δ and logarithmic above.

form filter needs a total time an order of magnitude larger than the others, as its Δ^{-2} scaling predicts. The optimized filter is shortest; its nonzero phases lower p_g (0.58 against 0.81 here), which the cost criterion accounts for. Second order cuts the step count by one to two orders of magnitude at these sizes. Figure 3 shows the three filter functions.

How much of the remaining gap could a better bound recover? For $N \leq 12$ we bisect a common scale factor on the certified k_i until the simulated infidelity just meets ϵ. The certified step count is 8.7–39× (median 21) larger than this oracle value (Table III). The oracle uses the exact ground state and is not a usable rule; it measures the price of certification at second order.

IV. GAP AND OVERLAP AS INDEPENDENT INPUTS

The cost of the filter depends on two physical inputs, the gap above the ground state and the trial state's weight on it, and on two choices, the size N and the target ϵ. Along any natural family of chains these move together: at fixed J_2 the gap closes with N, and at fixed bond dimension the MPS overlap decays with N. Varying J_2 does not help, because near the dimerized point a larger gap and a better MPS come together. We therefore construct a family in which Δ and γ are set independently at every N, and measure the filter on it.

A. Construction

Fixed gap. For each $N \in \{8, 10, 12, 14\}$ and target gap $g = E_1 - E_0 \in \{0.30, 0.35, \ldots, 0.55\}$ we find J_2 by bisection on exact diagonalization so that the singlet–triplet gap equals g to 10^{-9}. The gap closes with N at fixed J_2, so larger chains need larger J_2; pairs without a solution below the dimerized regime are left out (no J_2 gives $g \leq 0.35$ at $N = 8$ or $g = 0.30$ at $N = 10$). The tuned couplings reach $J_2 = 0.68$, beyond the Majumdar–Ghosh point, where a singlet could in principle drop below the triplet. For all 26 tuned (N, J_2), including $N = 16$, the lowest excitation in the $S^z = 0$ block, which contains a member of every multiplet, coincides with the lowest $S^z = 1$ state to 10^{-7}, so the first excitation is the triplet and g is the true gap (`results/gapov_e1check.json`).

Fixed overlap with a real circuit. The trial is the variational $\chi = 16$ MPS, prepared by its staircase circuit, followed by $R_y(\theta) = e^{-i\theta Y/2}$ on every even site. A uniform rotation of all sites would leave the overlap unchanged, because the ground state is a singlet; rotating one sublattice breaks that symmetry and lowers γ smoothly. θ is set by bisection so that γ hits the target to 10^{-12}. The trial cost is the MPS circuit plus $N/2$ single-qubit rotations at every γ, so the trial circuit does not change along an overlap scan.

Grid and measurements. Study A fixes $g \in \{0.40, 0.50\}$ and scans $\gamma \in \{0.03, 0.1, 0.3, 0.5, 0.7, 0.9\}$; study B fixes $\gamma \in \{0.1, 0.5\}$ and scans g over 0.30–0.55; both at $\epsilon \in \{10^{-2}, 10^{-3}, 10^{-4}\}$, 222 cases in all. Each case is run with both SBC schedules, dyadic and optimized, and with QETU for comparison (Sec. VII), and for each we record two costs. *Measured*, dyadic: the exact filter is applied to the trial state for each ladder (K, r); for the two shortest ladders that reach ϵ, the coarsest common Strang step at which the full Trotterized circuit still reaches ϵ is found by bisection on the state vector, and the cheaper is kept. *Measured*, optimized: SBC's design (Sec. II) for the target leakage ϵ with the true gap and overlap, then the same step search. The numerical design occasionally converges to a filter that barely passes the ground state, $f(E_0)^2 \sim 10^{-3}$, and with it a success probability far below γ; wherever $p < \gamma/2$ the design is repeated over four optimizer seeds and two design targets $(\epsilon, \epsilon/2)$ and the cheapest passing circuit is kept (21 of 286 points). *Certified*: Theorem 1 with exact energies (error bars 10^{-9}), the Weyl bound and the true γ. Both are reported per shot (one run of the circuit) and per success (divided by the success probability p), and, unless stated otherwise, for the filter alone, without the trial circuit.

B. The gap sets the circuit, the overlap sets the repetitions

Figures 4 and 5 show CNOTs, R_z rotations and evolution time per shot, and CNOTs per success, along both scans. The two inputs act on different resources.

Per shot, the overlap enters only logarithmically. At fixed gap, lowering γ thirtyfold (0.9 to 0.03) raises the measured CNOTs per shot of the dyadic filter by only 1.2–3.0× over all N and ϵ, and its evolution time by 1.5–3.0×. The filter must suppress the excited weight to $\eta^2 = \gamma\epsilon/(1 - \gamma)$; a worse trial raises $(1 - \gamma)/\gamma$ by

TABLE II. Plans for $N = 12$, $J_2 = 0.2411$, $\chi = 2$ trial. $T_{\rm phys} = T/W$ in units of $1/J_1$; CNOT and T counts per attempt (filter) and expected per success (trial + filter, divided by the certified success probability $(\sqrt{p_g} - \epsilon_U)^2$). "rot.": arbitrary rotations per attempt (trial + filter), each synthesized to accuracy δ. "meas.": simulated infidelity of the planned circuit. Source: `results/validate.json`.

ϵ	schedule	order	m	$T_{\rm phys}$	p_g	steps	CNOT (filter)	CNOT/success	rot.	δ	T/success	meas. $1-F$
10^{-2}	uniform	1	437	164.2	0.81	4,136,205	5.2×10^{8}	7.2×10^{8}	2.6×10^{8}	2×10^{-11}	3.8×10^{10}	–
10^{-2}	uniform	2	371	139.4	0.81	17,437	4.2×10^{6}	5.5×10^{6}	2.1×10^{6}	1×10^{-9}	2.4×10^{8}	1.5×10^{-3}
10^{-2}	dyadic	1	10	23.3	0.81	77,308	9.7×10^{6}	1.4×10^{7}	4.9×10^{6}	1×10^{-9}	6.1×10^{8}	–
10^{-2}	dyadic	2	10	23.3	0.81	884	2.1×10^{5}	2.9×10^{5}	1.1×10^{5}	5×10^{-8}	1.1×10^{7}	8.5×10^{-7}
10^{-2}	optimized	1	6	8.7	0.58	13,384	1.7×10^{6}	3.2×10^{6}	8.4×10^{5}	5×10^{-9}	1.3×10^{8}	1.1×10^{-3}
10^{-2}	optimized	2	6	8.7	0.58	226	5.4×10^{4}	1.0×10^{5}	2.7×10^{4}	2×10^{-7}	3.5×10^{6}	1.1×10^{-3}
10^{-3}	uniform	1	745	279.9	0.81	27,817,555	3.5×10^{9}	4.5×10^{9}	1.8×10^{9}	1×10^{-12}	2.7×10^{11}	–
10^{-3}	uniform	2	658	247.2	0.81	55,930	1.3×10^{7}	1.7×10^{7}	6.7×10^{6}	2×10^{-10}	8.3×10^{8}	–
10^{-3}	dyadic	1	15	34.9	0.81	536,160	6.8×10^{7}	8.7×10^{7}	3.4×10^{7}	5×10^{-11}	4.5×10^{9}	–
10^{-3}	dyadic	2	15	34.9	0.81	2,835	6.8×10^{5}	8.8×10^{5}	3.4×10^{5}	5×10^{-9}	3.6×10^{7}	2.5×10^{-9}
10^{-3}	optimized	1	8	14.4	0.67	80,421	1.0×10^{7}	1.6×10^{7}	5.1×10^{6}	4×10^{-10}	7.4×10^{8}	–
10^{-3}	optimized	2	8	14.4	0.67	709	1.7×10^{5}	2.6×10^{5}	8.5×10^{4}	2×10^{-8}	1.0×10^{7}	6.2×10^{-5}

TABLE III. Certified over oracle step count (second order, $\chi = 2$). Columns: $N = 8, 10, 12$ for each J_2. Source: `results/oracle.json`.

		$J_2 = 0$			$J_2 = 0.2411$			$J_2 = 0.4$		
ϵ	schedule	8	10	12	8	10	12	8	10	12
1×10^{-2}	dyadic	13	22	30	24	24	34	18	15	28
1×10^{-2}	optimized	9	12	13	14	16	21	11	12	17
1×10^{-3}	dyadic	22	24	39	31	36	39	20	27	28
1×10^{-3}	optimized	12	17	21	24	22	27	14	16	17

a factor of 290 here, which adds only $\ln 290 \approx 5.7$ to $L = \ln[(1-\gamma)/(\gamma\epsilon)]$, already 2–13 across these cases. The dyadic filter grows in whole pulses, so a worse trial often needs no extra pulse at all; its per-shot cost moves in steps (Fig. 4, top).

Per shot, the gap enters as a power. At fixed overlap, closing the gap by 1.8× (0.55 to 0.30) raises the CNOTs per shot by 2.7–4.3× and the evolution time by 1.5–2.8× (Fig. 5). The time grows about as $1/\Delta$, the cost of resolving a smaller gap. Both behaviors are the textbook $T = O(\Delta^{-1}\ln\eta^{-1})$ of filter-based preparation [6, 7]; what the controlled study adds is that the measured circuits follow it at fixed N with the two inputs varied separately, and how the gates per shot and the Trotter step add to it.

Per success, the overlap returns as $1/\gamma$. Every run succeeds with probability $p \approx \gamma f(E_0)^2$, and for the dyadic filter $f(E_0) = 1$, so the expected cost per success carries the factor $1/\gamma$ (last columns of Figs. 4 and 5). This is the only place where the overlap enters as a power. The practical reading is direct: the gap decides the depth of the coherent circuit, and with it whether the filter fits within a coherence or error budget at all; the overlap decides how many times that circuit has to be run.

R_z counts follow CNOT counts exactly. Every controlled bond block has 3 R_z and 6 CNOTs, and each pulse adds one ancilla rotation, so $R_z/\mathrm{CNOT} = 0.50$ in every case. Every statement about CNOTs below therefore holds for the arbitrary rotations, and so for the T count of a fault-tolerant implementation up to the synthesis factor $3\log_2(1/\delta)$.

C. Mechanism: filter length, Trotter step and gates per step

The per-shot cost factorizes as

$$\text{CNOTs per shot} = \frac{T}{\delta t} \times c_{\rm step}(N), \tag{11}$$

and each factor answers to different inputs (Fig. 6).

Filter length. The physical evolution time T (in $1/J_1$) depends on Δ and L only. Fitting all cases to $a\,N^{\alpha}\Delta^{-\beta}L^{\lambda}\epsilon^{-e}$, the measured dyadic filter has $T \propto \Delta^{-1.19}L^{0.77}$ with $N^{-0.24}$, and the certified one $T \propto \Delta^{-1.10}L^{0.84}$ with $N^{0.04}$: at a fixed gap in absolute units, the filter's length does not depend on the size of the system. The certified length is a median 2.0× the measured one (range 1.2–3): the certificate sizes the filter nearly right.

Trotter step. The measured dyadic filter tolerates very coarse steps: median $\delta t = 1.25$, 0.56, 0.30 (in $1/J_1$) at $\epsilon = 10^{-2}$, 10^{-3}, 10^{-4}, growing with the gap. Its pulses are long and smooth, and the filtered state lives at the

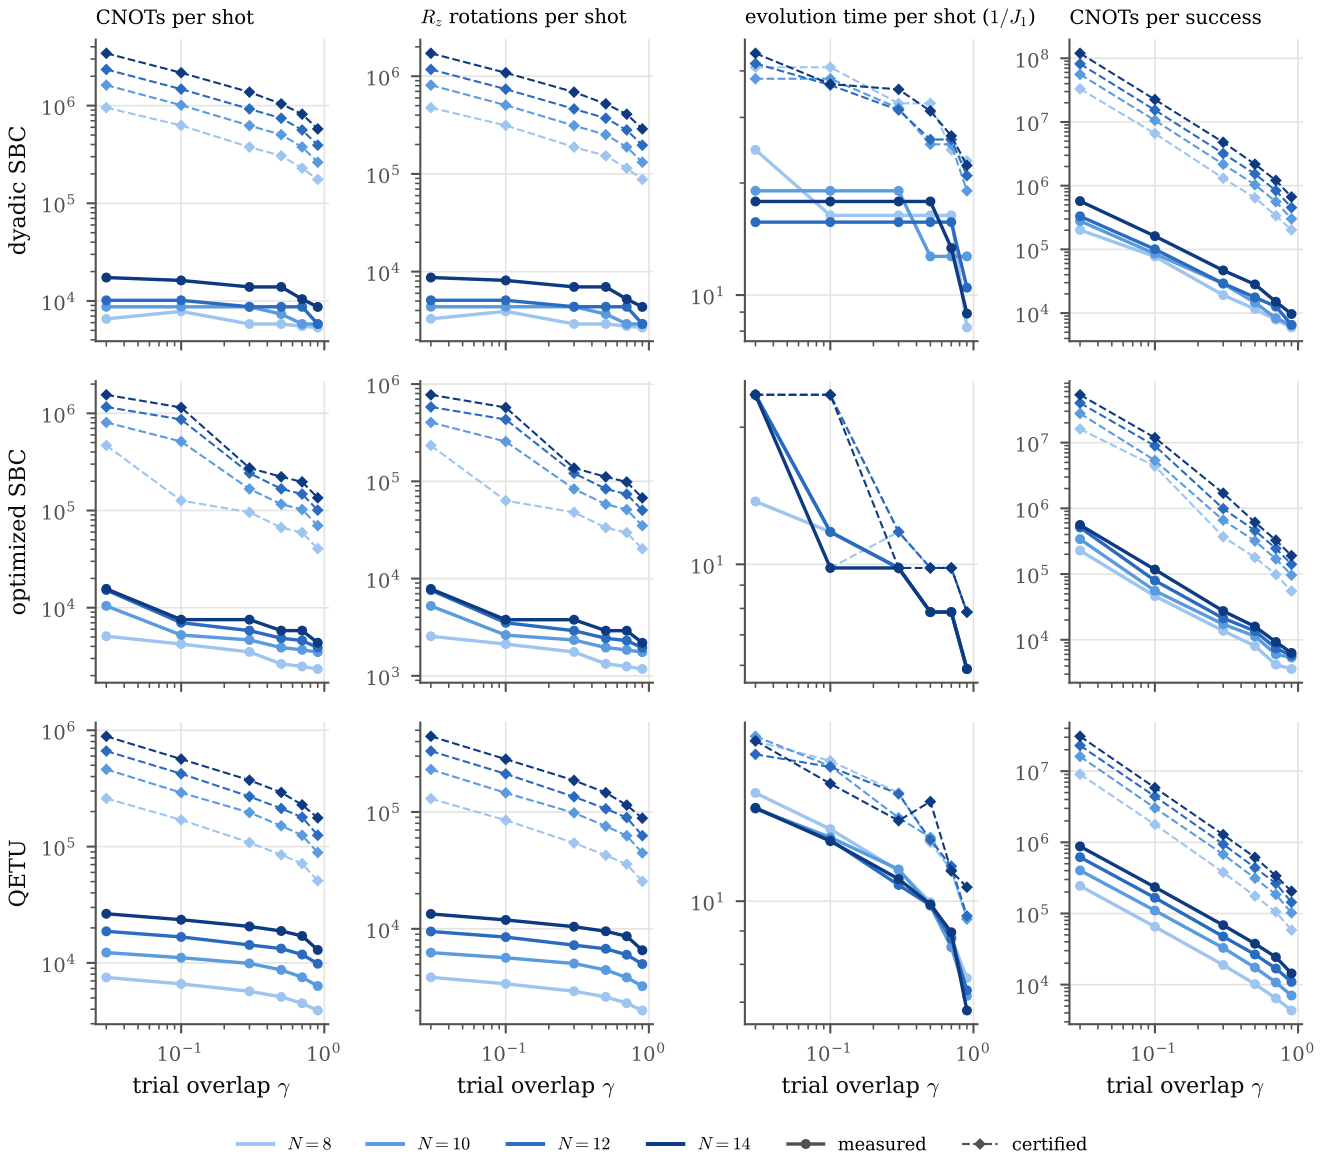


FIG. 4. Study A: filter-only resources at a gap held at $g = 0.40$ for every N, $\epsilon = 10^{-3}$, against the trial overlap. Rows: dyadic SBC, optimized SBC, QETU (Sec. VII). Solid: measured by simulation; dashed: certified. Columns 1–3 per shot, column 4 per success. Source: `results/gapov.jsonl`, `results/gapov_filter.json`.

bottom of the spectrum, where Trotter errors are much smaller than their worst case [19]. The certified step must cover the worst case over the whole evolution: median $\delta t = 0.017$, 0.009, 0.005, about 60–75 times finer, and shrinking as the total time grows. As a result the certified step count grows as $L^{2.03}$ against $L^{0.71}$ measured, and is a median 89.4× (range 26.2–339) the measured one. Almost all of the certificate's overhead is in the Trotter bound, not in the filter.

Gates per step. One Strang step costs a median $20\,N$ CNOTs for the two-neighbor chain ($6(2B-2)$ with $B = 2N-3$ bonds). The measured dyadic step count barely grows with N (exponent 0.08), so its CNOTs grow as $N^{1.30}$; the certified step count grows as $N^{0.85}$ through the commutator constant and the scaled time, giving $N^{2.08}$.

D. Cost laws

Fitting the expected CNOTs per success (filter and trial circuit, as a user would pay them) to

$$C \approx a\, N^{\alpha}\, \Delta^{-\beta}\, L^{\lambda}\, \gamma^{-g}, \qquad L = \ln \frac{1-\gamma}{\gamma\epsilon}, \tag{12}$$

over all 222 cases gives Table IV. For the dyadic filter,

$$C_{\text{meas}} \approx 10.1\, N^{1.50} \Delta^{-1.15} L^{1.41} \gamma^{-0.78}, \tag{13}$$

$$C_{\text{cert}} \approx 15\, N^{2.07} \Delta^{-1.27} L^{2.43} \gamma^{-0.98}, \tag{14}$$

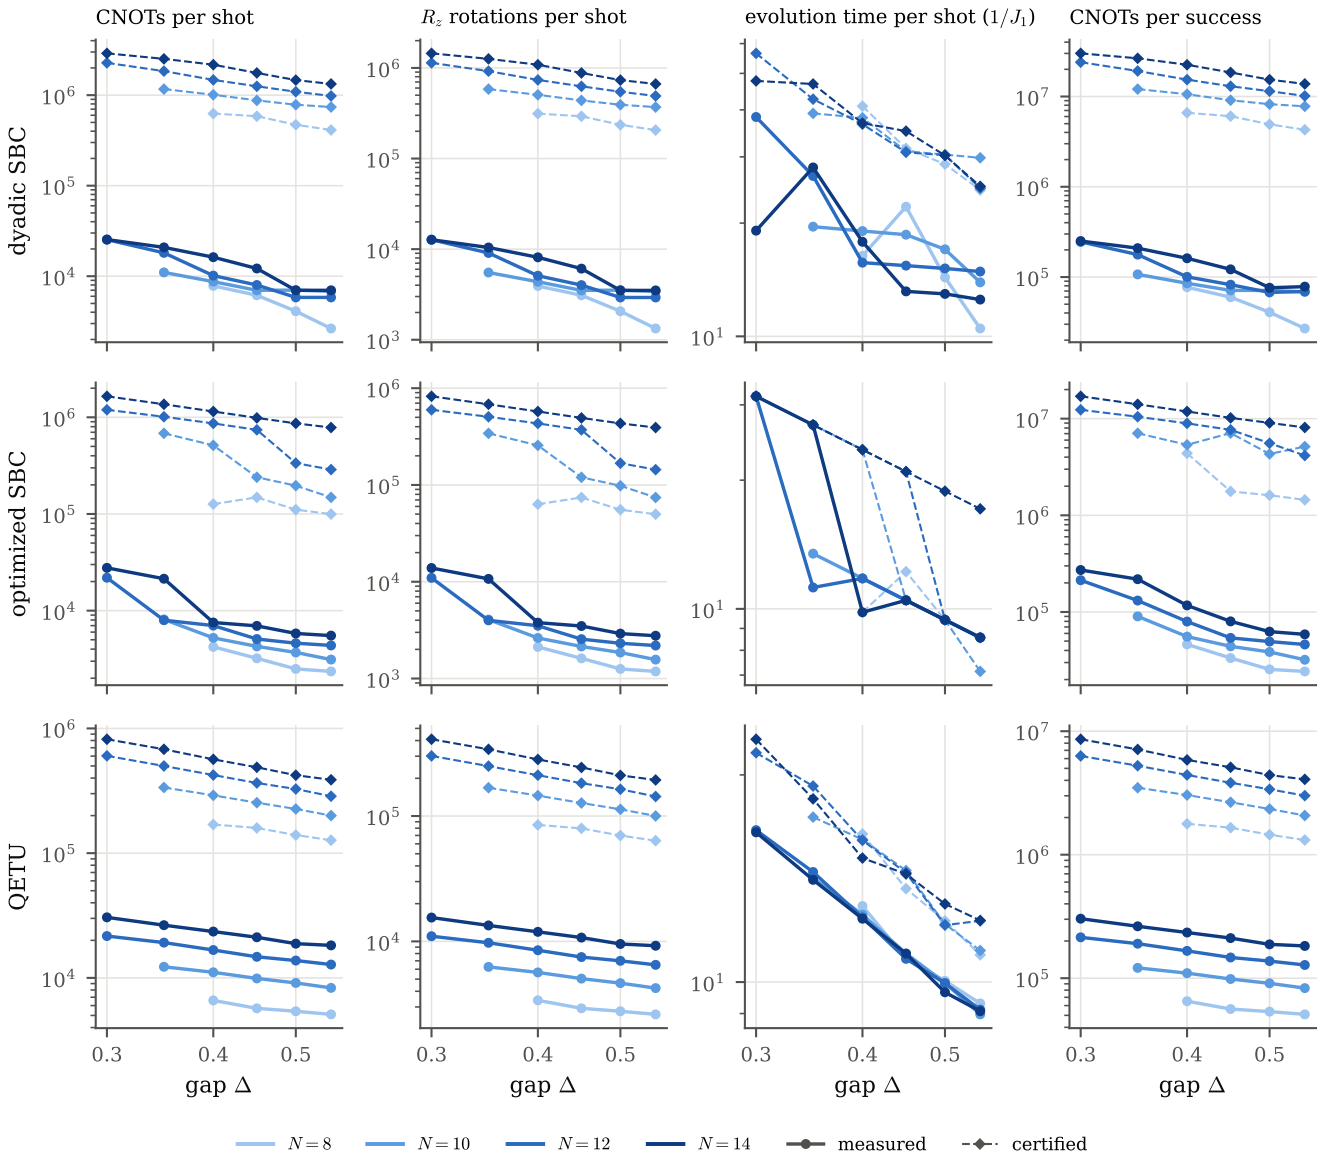


FIG. 5. Study B: as Fig. 4 at an overlap held at $\gamma = 0.10$ for every N, against the gap (held fixed across N by tuning J_2).

within a typical factor of 1.23 and 1.22 respectively. The certified law has the $1/\gamma$ of the repetitions and an $L^{2.4}$ from the Trotter bound; the measured law grows more slowly in both N and L. The measured dyadic γ exponent is below one because its success probability ranges from 0.76 to 1.24 times γ: at the coarse steps it tolerates, each Trotterized pulse is slightly off-diagonal in the eigenbasis and moves weight into or out of the ground state. Part of its γ dependence is therefore absorbed into L. Across all cases the certified cost lies 14–258× (median 62) above the measured one. These exponents summarize $N = 8$–14 and $g = 0.30$–0.55; they describe the data and are not asymptotic laws.

Appendix B compares the two SBC schedules case by case and maps the cost over the (Δ, γ) plane.

V. RESOURCES AND REGIMES OF ADVANTAGE

The controlled study fixes the inputs; here they take the values that DMRG supplies for natural chains, at sizes beyond exact simulation, and we ask when filtering beats the alternative of compiling a better MPS.

A. Resources up to $N = 48$

Without simulation the pipeline only needs DMRG, so we plan for N up to 48 with the same $\chi = 2$ trial family (Fig. 7, Table V). Over all sizes, couplings and targets the second-order plan is 17–666× (median 93) cheaper

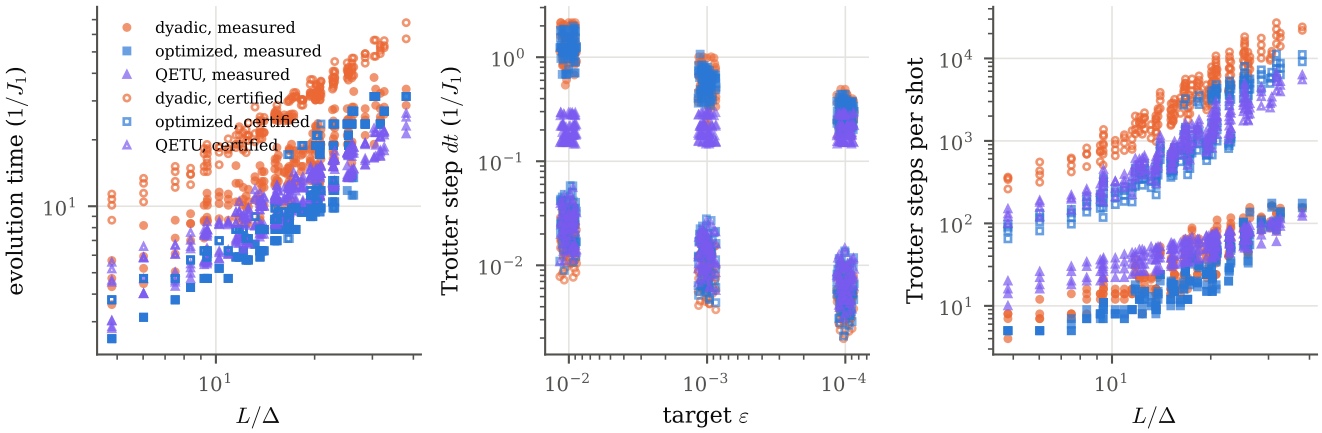


FIG. 6. All 222 cases. (a) Evolution time against L/Δ, $L = \ln[(1-\gamma)/(\gamma\epsilon)]$. (b) Trotter step against ϵ (jittered horizontally). (c) Trotter steps per shot against L/Δ. Filled: measured; open: certified. Orange circles: dyadic SBC; blue squares: optimized SBC; purple triangles: QETU.

TABLE IV. Fits of the expected CNOTs per success to Eq. (12) over all 222 gap/overlap cases. "rms"/"max": typical and worst factor between a case and its fit. Source: `results/gapov.jsonl`.

	cases	a	α	β	λ	g	rms	max
dyadic, measured	222	1.0×10^1	1.50	1.15	1.41	0.78	1.23	1.9
optimized, measured	222	9.8	1.50	1.36	1.24	0.83	1.30	2.3
QETU, measured	222	1.3×10^1	2.09	0.74	0.68	0.99	1.04	1.2
dyadic, certified	222	1.5×10^1	2.07	1.27	2.43	0.98	1.22	2.2
optimized, certified	222	1.9	2.27	1.30	2.74	0.97	1.45	3.2
QETU, certified	222	9.5	2.03	1.11	2.17	0.99	1.22	2.0

than the first-order plan of the same schedule. The optimized schedule is 0.7–3.7× (median 2.9) cheaper than the dyadic one at second order, and more expensive in 1 of 29 comparable cases. Its numerical design, however, failed in 14 of 144 plans, all at $N \geq 32$ where $\gamma \lesssim 0.4$ or $\epsilon = 10^{-3}$ demand a suppression η the SLSQP search with $F(0)^2 \geq 10^{-3}$ did not reach; the dyadic schedule is feasible everywhere, which is why the pipeline offers an automatic choice between the two. For $N = 16$–48 the expected CNOT count of the dyadic second-order plan grows as $N^{4.9}$, $N^{5.0}$ and $N^{3.5}$ for $J_2 = 0, 0.2411, 0.4$, and its step count as $N^{2.6}$, $N^{2.7}$, $N^{2.1}$. These exponents combine the shrinking gap ($T \propto 1/\Delta$, and $\Delta = g/W$ with $W \propto N$), the growing commutator constant, the number of terms per step and the decaying overlap (through the success probability); they are fits over $N = 16$–48, not asymptotic statements. At $N = 48$, $J_2 = 0.2411$ the best $\chi = 2$ plan (dyadic) needs 2.7×10^8 CNOTs per success, which Sec. V C reduces by choosing a better trial.

B. What if the energy estimates are wrong?

The certificate is only as good as its inputs, so we perturb them for $N = 12$, $J_2 = 0.2411$, $\epsilon = 10^{-2}$ (optimized, second order) and simulate the resulting circuit on the true trial state (Fig. 8).

Gap. A gap assumed larger than the true one moves the stop-band edge above the first excitation, which then sits in the filter's transition region. The target was missed in 3 cases, $f \geq 1.5$, with infidelity up to 5.1×10^{-2} at $f = 3$; at $f = 1.1$ and 1.25 the looseness of the certificate still covered the error. Underestimating the gap is safe and cheap: even $f = 0.5$ costs only 2.9×. Since DMRG gives an *upper* bound on E_1, an explicit margin δE_1 is the natural safeguard.

Ground energy. An error bar δE_0 widens the pass band and lowers F_{pass}; the cost stays flat until δE_0 is a sizeable fraction of the gap (13.5× at $\delta E_0 = 0.1$ gap). DMRG energies are far more accurate than this, so E_0 is never the limiting input here.

Overlap. Underestimating γ is safe (a stricter η); the planned cost rises 9.2× at $f = 0.25$, mostly because the planned success probability $p_g = \gamma F_{\text{pass}}^2$ uses the pes-

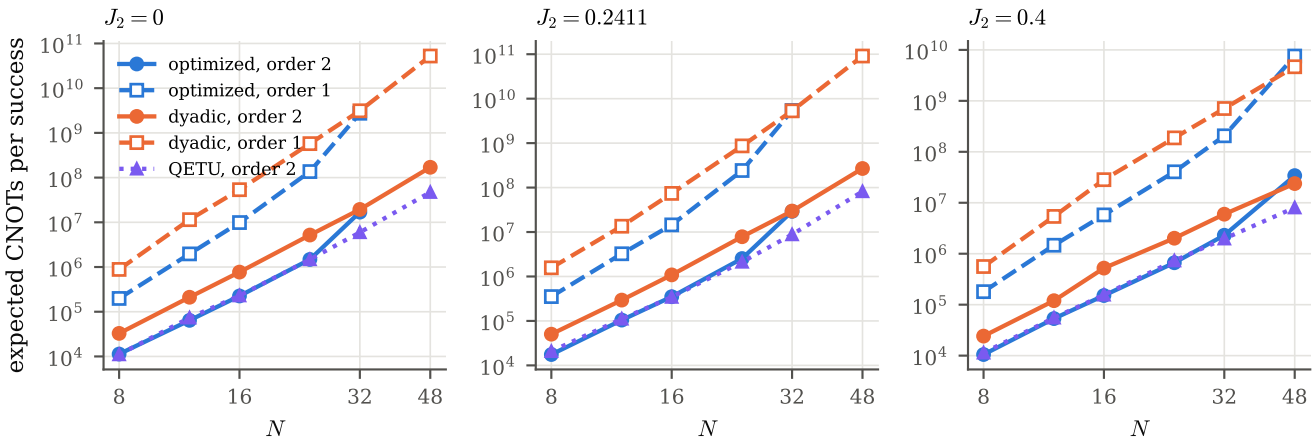


FIG. 7. Expected CNOTs per successful preparation versus N (trial $\chi = 2$, $\epsilon = 10^{-2}$), including the trial circuit and the repetitions (certified success probability). Solid: second order; dashed: first order. Missing optimized points: no feasible design (see text). Dotted triangles: QETU, second order. Sources: `results/scaling.json`, `results/ext_qetu_scaling.json`.

TABLE V. Cheapest second-order plan ("sch.": dyadic or optimized) for $\epsilon = 10^{-2}$, trial $\chi = 2$. Δ is the scaled gap, $T_{\text{phys}} = T/W$ in $1/J_1$. Steps for first (1st) and second (2nd) order with the same schedule; CNOT and T counts for second order per attempt (att.) and per success (succ.), divided by the certified success probability; "rot.": arbitrary rotations per attempt, synthesized to accuracy δ. Source: `results/scaling.json`.

N	J_2	γ	Δ	sch.	T_{phys}	p_g	steps (1st)	steps (2nd)	CNOT/att.	CNOT/succ.	rot.	δ	T/succ.
16	0	0.475	0.0206	opt.	14.3	0.38	38,020	464	7.8×10^4	2.3×10^5	3.9×10^4	9×10^{-8}	8.0×10^6
16	0.2411	0.703	0.0192	opt.	15.0	0.58	41,572	521	1.8×10^5	3.5×10^5	8.8×10^4	7×10^{-8}	1.3×10^7
16	0.4	0.911	0.0211	opt.	6.6	0.63	19,266	264	8.9×10^4	1.5×10^5	4.5×10^4	6×10^{-8}	5.4×10^6
24	0	0.328	0.0094	opt.	20.5	0.27	251,160	1,422	3.8×10^5	1.5×10^6	1.9×10^5	8×10^{-9}	6.0×10^7
24	0.2411	0.517	0.0087	opt.	16.3	0.29	248,256	1,329	7.0×10^5	2.5×10^6	3.5×10^5	5×10^{-9}	1.1×10^8
24	0.4	0.839	0.0102	opt.	10.8	0.59	87,533	700	3.7×10^5	6.6×10^5	1.9×10^5	1×10^{-8}	2.6×10^7
32	0	0.234	0.0054	opt.	80.0	0.23	3,103,919	9,847	3.5×10^6	1.7×10^7	1.8×10^6	1×10^{-9}	7.4×10^8
32	0.2411	0.373	0.0049	opt.	85.8	0.37	4,979,405	13,617	9.8×10^6	2.9×10^7	4.9×10^6	6×10^{-10}	1.3×10^9
32	0.4	0.768	0.0062	opt.	13.3	0.41	221,517	1,228	8.8×10^5	2.3×10^6	4.4×10^5	5×10^{-9}	9.6×10^7
48	0	0.125	0.0025	dya.	146.0	0.13	20,734,266	34,380	1.9×10^7	1.7×10^8	9.5×10^6	2×10^{-10}	8.2×10^9
48	0.2411	0.191	0.0022	dya.	142.5	0.19	27,888,279	41,586	4.6×10^7	2.7×10^8	2.3×10^7	1×10^{-10}	1.3×10^{10}
48	0.4	0.640	0.0031	dya.	69.3	0.64	4,816,485	12,414	1.4×10^7	2.4×10^7	6.9×10^6	6×10^{-10}	1.1×10^9

simistic γ, while the actual success rate does not change. Overestimating it by 2% did not break the target in this instance (0 failures), because the leakage term is itself loose.

All instances. To see whether this one instance is typical we repeated the gap and overlap perturbations for every $N \leq 14$, all three couplings and the dyadic, optimized and QETU schedules (second order, $\epsilon = 10^{-2}$; Table VI). With exact inputs all 36 plans meet the target. An overestimate of the gap by 10% already breaks it in 2 of 36 plans (marginally, worst 1.1×10^{-2}), by 25% in 10 and by 50% in 16. The closed-form dyadic schedule is the most tolerant (0 misses at 25%), because its bound is the loosest; the optimized and Chebyshev filters, whose leakage terms are nearly tight, have no slack to absorb a wrong gap. Overestimating γ broke the target in 12 of 72 plans, all of them through the same mechanism: the inflated γ crossed $1 - \epsilon$, so the pipeline dropped the filter altogether (12 of 12). A margin on both inputs is therefore not optional: δE_1 should be a sizeable fraction of the gap when E_1 is not certified, and the empty filter should be accepted only on a certified (for example Eckart) lower bound on γ.

A margin from DMRG convergence. A rule that uses only DMRG data is to set δE_0 and δE_1 to the change of E_0 and E_1 between the last two bond dimensions of the DMRG ladder, which is the pipeline's default error bar. We tested it on 15 systems with exact energies (chains

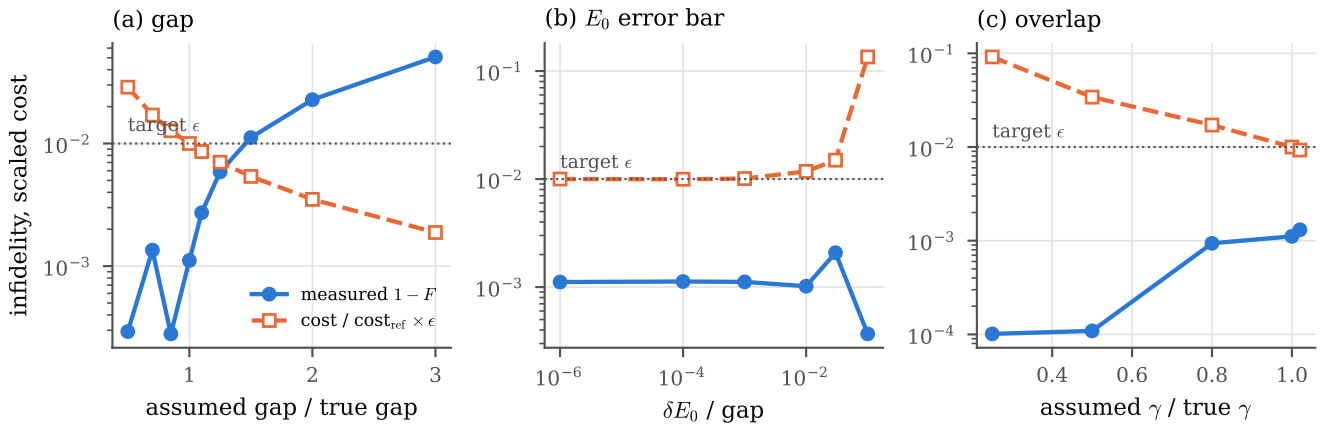


FIG. 8. Sensitivity to the inputs ($N = 12$, $J_2 = 0.2411$, $\epsilon = 10^{-2}$). Circles: measured infidelity of the planned circuit on the true trial state. Squares: expected CNOT cost relative to the plan with exact inputs, multiplied by ϵ to share the axis. (a) Assumed gap $f\times$ the true gap; (b) error bar on E_0; (c) assumed overlap $f\times$ the true γ. Source: `results/sensitivity.json`.

$N = 12$–16, cylinders 4×3–4×5, all J_2), stopping the ladder at each of $\chi = 8$, 16, 32, 64, 128 (75 cases; `run_margin.py`). Without error bars the DMRG energies are consistent with the exact ones in only 29 of 75 cases (the DMRG gap is up to 17% too large); with them the intervals contain both exact energies in 75 of 75. On the four cylinders with full spectra we then planned every schedule at $\epsilon = 10^{-2}$–10^{-4} from each stopped ladder and applied the exact filter function to the trial's spectral weights: all 396 plans met the target. Once DMRG has converged ($\chi \geq 64$) the margin costs a median 1.00× (at most 1.4×); at $\chi = 32$ a median 1.6×; at $\chi \leq 16$ it makes the energy window inconsistent in 58 of 72 plans, which the pipeline reports instead of returning a plan. The rule therefore either covers the gap at small cost or says that DMRG is not converged enough to plan from. In these instances the plans without margin also passed, because their gap errors at $\chi \geq 8$ were below the 10% at which Table VI starts to fail; the test shows that the margin is valid and cheap, not that it was needed.

C. Choosing the ansatz

The trial bond dimension trades circuit depth for overlap. Figure 9 and Table VII show the expected total cost for $\chi = 1$ to 8 (the pipeline choosing the cheaper of the dyadic and optimized schedules). A product state ($\chi = 1$) has an exponentially small overlap and is by far the most expensive start (10^7–10^{23} CNOTs). Going to $\chi = 2$ and $\chi = 4$ lowers the cost by orders of magnitude: at $N = 48$, $J_2 = 0.2411$ the $\chi = 4$ plan needs 4.8×10^6 CNOTs and 2.1×10^8 T gates per success. At $\chi = 8$ the trial state alone satisfies $1 - \gamma \leq 10^{-2}$ in 9 of 9 cases ($\gamma = 0.9946$ at $N = 48$, $J_2 = 0.2411$), and the pipeline returns the empty filter: 4,500 CNOTs. For this gapless-to-weakly-gapped one-dimensional chain at $\epsilon = 10^{-2}$, directly compiling a better MPS is therefore much cheaper than filtering a worse one. The filter becomes the better tool when the target is tighter than what a compilable MPS reaches, when a certificate on the output is required, or when the classical method cannot provide a compact high-overlap trial circuit, as for two-dimensional lattices.

D. Measured crossover

The comparison above is certified on both sides only for the MPS. To compare like with like we measure the filtered route too: for chains of $N = 12$–16 at $J_2 = 0$, 0.2411 and 0.4, the exact spectrum of each Hamiltonian is computed per S^z block, every MPS trial with $\chi = 2, \ldots, 64$ is decomposed on it, and for each target $\epsilon \in [10^{-5}, 10^{-2}]$ the cheapest route is found: the smallest MPS that meets $1 - \gamma \leq \epsilon$ alone, against the cheapest smaller MPS followed by a measured dyadic or optimized SBC filter (shortest exact ladder or design, then the coarsest passing Strang step), with cost (trial + filter)/p.

Figure 10 shows the result on nine chains. Compiling the MPS directly is the cheaper route in most cases: the filtered route wins in 18 of the 63 (chain, ϵ) targets, by 1.1–3.2×. The wins share a single mechanism. Because MPS compilation cost jumps with χ (here by a factor of about four from $\chi = 4$ to 8 and from 8 to 16), a target lying just below what a given χ reaches forces the MPS-alone route up to the next bond dimension, while the trial at the lower χ misses the target only slightly. All 18 wins filter such a trial, with $(1-\gamma)/\epsilon \leq 8.8$ (13 of them below 3), using a dyadic ladder of one to four pulses ($K \leq 3$, $r \leq 2$, 1–21 Strang steps) that succeeds with $p \geq 0.988$. Counting every candidate rather than only the cheapest per target (a *filtered route* is one trial χ followed by its

TABLE VI. Gap and overlap overestimation for every $N \in \{8, 10, 12, 14\}$ and J_2 (12 instances), $\epsilon = 10^{-2}$, second order: number of plans that miss the target and the worst measured infidelity. f: assumed over true gap or γ. Source: `results/ext_sensitivity.json`.

		dyadic		optimized		QETU	
input	f	miss	worst	miss	worst	miss	worst
gap	1	0/12	1.7×10^{-3}	0/12	4.2×10^{-3}	0/12	2.3×10^{-3}
gap	1.1	0/12	6.5×10^{-3}	1/12	1.1×10^{-2}	1/12	1.0×10^{-2}
gap	1.25	0/12	6.8×10^{-3}	5/12	2.6×10^{-2}	5/12	3.0×10^{-2}
gap	1.5	2/12	6.3×10^{-2}	7/12	6.8×10^{-2}	7/12	9.2×10^{-2}
gap	2	8/12	1.4×10^{-1}	7/12	1.5×10^{-1}	11/12	2.2×10^{-1}
gap	3	11/12	2.2×10^{-1}	12/12	2.5×10^{-1}	12/12	3.1×10^{-1}
γ	1.02	1/12	2.4×10^{-2}	1/12	2.4×10^{-2}	1/12	2.4×10^{-2}
γ	1.05	3/12	5.5×10^{-2}	3/12	5.5×10^{-2}	3/12	5.5×10^{-2}

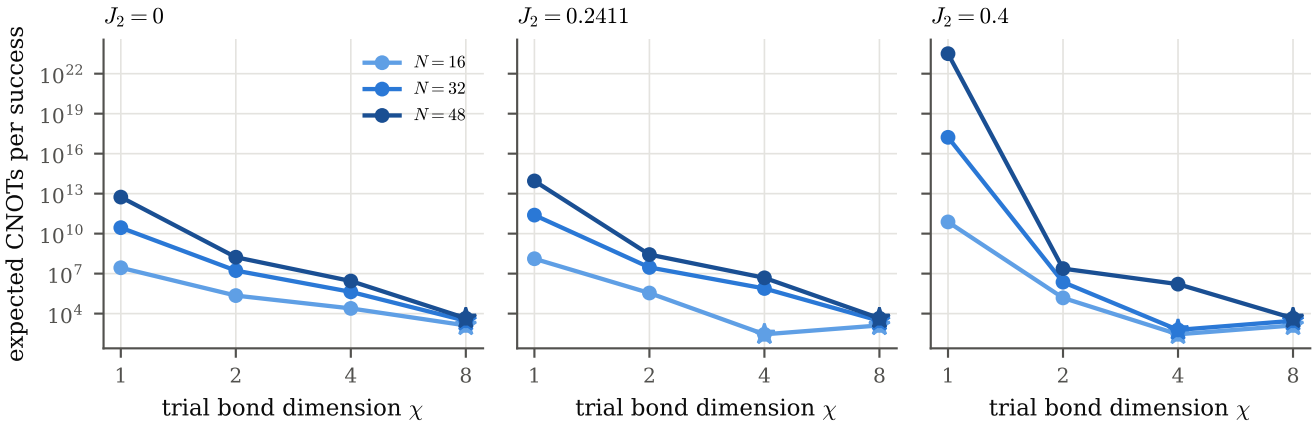


FIG. 9. Expected CNOTs per success versus trial bond dimension ($\epsilon = 10^{-2}$, second order, cheaper of dyadic and optimized). Stars: the trial state alone satisfies $1 - \gamma \leq \epsilon$ and no filter is used. Source: `results/ansatz.json`.

measured filter, at one target; each target has one per χ below the one that meets it alone, 121 in all), of the 31 filtered routes with $(1-\gamma)/\epsilon \leq 8.8$, 18 beat the MPS alone; all 90 with a larger ratio lose, by a median factor of 20. In this regime the filter acts as a last-mile polisher: it removes the residual infidelity of a nearly sufficient trial, but recovering a larger factor is cheaper by doubling χ. The location of the wins in ϵ therefore depends on the system through where each χ lands, not on N or J_2 in any simple way, and the wins are an artifact of the bond-dimension grid as much as a property of the filter: Sec. VI gives the MPS alone every $\chi \leq 64$, and they disappear. The optimized design never wins. A single dyadic pulse lasts only $\pi/(1+\Delta)$ and keeps $f(E_0) = 1$; it does not resolve the gap, but it removes 17–68% of the residual weight, which suffices for a trial this close to the target. The optimized design certifies suppression on the whole band above Δ and so needs $T \gtrsim \pi/\Delta$: rerunning the competitive cases ($(1-\gamma)/\epsilon \leq 20$; 32 of 40 completed) with one- to three-pulse designs and the design target loosened up to 10ϵ lowered its cost by up to 3× but left $T = 35$–66 against 3–21 for the winning dyadic ladders, and it beat the dyadic route in 3 of the 23 cases with $(1-\gamma)/\epsilon \leq 8.8$ and the MPS alone in none (`results/emp_actual1d_opt.jsonl`). QETU beats the MPS alone in 11 of the targets but is never the cheapest route. This matches the observation for rodeo filters that a ratio-2 time series is optimal when the total time is short and longer series favor other designs [23]. Because the bond dimension a chain needs for fixed ϵ grows only slowly with N, we expect this picture to hold for longer chains; on cylinders it reverses (Sec. VI).

VI. TWO DIMENSIONS: WHERE THE FILTER PAYS

The pipeline is generic in H, so we ran the measured crossover of Sec. V D on the J_1–J_2 model on 4×3 and 4×4 cylinders (periodic around the four-site circumference,

TABLE VII. Overlap γ and expected CNOTs per success for trial bond dimension $\chi = 1, 2, 4, 8$ ($\epsilon = 10^{-2}$, second order, cheaper schedule; no filter when $1 - \gamma \leq \epsilon$).

N	J_2	γ_1	CNOT	γ_2	CNOT	γ_4	CNOT	γ_8	CNOT
16	0	0.028	2.8×10^7	0.475	2.3×10^5	0.984	2.5×10^4	1.000	1.3×10^3
16	0.2411	0.019	1.3×10^8	0.703	3.5×10^5	0.998	2.8×10^2	1.000	1.3×10^3
16	0.4	0.000	7.6×10^{10}	0.911	1.5×10^5	0.999	2.8×10^2	1.000	1.3×10^3
32	0	0.001	2.8×10^{10}	0.234	1.7×10^7	0.913	4.3×10^5	0.998	2.9×10^3
32	0.2411	0.001	2.5×10^{11}	0.373	2.9×10^7	0.972	7.7×10^5	0.999	2.9×10^3
32	0.4	0.000	1.7×10^{17}	0.768	2.3×10^6	0.990	6.0×10^2	1.000	2.9×10^3
48	0	0.000	5.5×10^{12}	0.125	1.7×10^8	0.817	2.8×10^6	0.990	4.5×10^3
48	0.2411	0.000	9.0×10^{13}	0.191	2.7×10^8	0.920	4.8×10^6	0.995	4.5×10^3
48	0.4	0.000	3.2×10^{23}	0.640	2.4×10^7	0.976	1.6×10^6	0.998	4.5×10^3

open along L_x) at $J_2 = 0$ and 0.3, with snake-path MPS trials up to $\chi = 64$ and exact spectra (Fig. 11). Cylinders are much harder for a snake MPS than chains. At 4×4, $J_2 = 0$, the $\chi = 16$ trial still misses $\epsilon = 10^{-2}$ ($1 - \gamma = 1.1 \times 10^{-2}$), $\chi = 32$ stalls at 7×10^{-4}, and no $\chi \leq 64$ trial reaches 10^{-5} ($1 - \gamma \approx 2 \times 10^{-5}$ at $\chi = 64$), while the compile cost grows about fourfold per doubling of χ (5 328, 20 548 and 76 600 CNOTs for $\chi = 16$, 32, 64). On 4×3, $\chi = 64$ is exact, but only at 45 960 CNOTs.

The result is the opposite of the chain. With the uniform MPS cost model, the filtered route is cheaper at 27 of the 28 cylinder targets from $\epsilon = 10^{-2}$ to 10^{-5}, by 1.1–8.8× where both routes exist, and it is the only MPS-based route at 10^{-5} on 4×4. The exception is $\epsilon = 10^{-2}$ on 4×4, $J_2 = 0.3$, where the $\chi = 16$ MPS alone already reaches 8.3×10^{-3}. Every win uses the dyadic ladder, on a $\chi = 16$ trial in 24 of the 27 (otherwise $\chi = 4$, 8 or 32), with $p \geq 0.77$ and 1–46 Strang steps. Unlike the chain, the filter here recovers large factors: the trials miss the target by $(1-\gamma)/\epsilon$ up to 10^3, and the advantage grows toward tight targets (up to 8.8× at 4×4, $J_2 = 0$, $\epsilon = 3 \times 10^{-4}$, where $\chi = 16$ plus 10 Strang steps replaces $\chi = 64$). Two properties of the filter carry this: its depth grows only with $\ln(1/\epsilon)$ at a fixed gap, while the MPS needed to reach ϵ alone grows by a factor of four per bond-dimension step, and its cost in the overlap is the logarithm of Sec. IV. On the 16-site chain in the same comparison the filter wins only at $\epsilon = 3 \times 10^{-5}$ and 10^{-5}.

Stronger baselines. The result depends on how the MPS circuits are costed, and the uniform staircase bound used so far is generous to the filter: it charges every site a unitary on $b+1$ qubits even where the bond is smaller, so at $\chi = 64$ it charges 45 960 and 76 600 CNOTs for states of 12 and 16 qubits that generic state preparation [38] reaches exactly with 3 831 and 62 423. Table VIII therefore repeats every crossover, measured and certified, with two further cost models that apply to both routes. The first is constructive: each site unitary is sized by its actual bonds, and exact generic state preparation is offered as an alternative route at every ϵ. The second is a floor that no compiler of these circuits can beat: the isometry lower bound [37] per site with the actual bonds, $\lceil (2^{m+n+1} - 4^m - 3n - 1)/4 \rceil$ CNOTs for an isometry from m to n qubits, and its $m = 0$ case for generic state preparation (2 039 and 32 756 CNOTs). On 4×4 the measured advantage survives both: the filtered route wins 13 of 14 targets (up to 5.7×) with the constructive model and 13 of 14 (up to 3.2×) against the floor. On 4×3 it does not: 12 qubits are few enough that exact generic preparation costs less than most filtered routes, and the wins fall to 5 and 3 of 14. On chains they fall from 18 to 16 and 8 of 63. The measured two-dimensional result is therefore a statement about 4×4, the largest cylinder we can simulate, and it is robust there because the filter's cost grows polynomially in N while generic preparation grows as 2^N and a snake MPS needs a bond dimension that grows with the width. Certified, the filtered route wins none of the 91 targets under either stronger model (median certified cost 8.5 and 18.9 times the best alternative on 4×4), so the "only route" statement above holds only among MPS-based routes. *Intermediate bond dimensions.* The filtered route is a minimum over trial χ, filter design and Trotter step, while the MPS alone has only $\chi \in \{2, 4, \ldots, 64\}$. To give the MPS alone the same freedom we added every crossover system's variational MPS at $\chi = 3$, 5, 6, 10, 12, 20, 24, 40 and 48 (`run_chi_mid.py`; Table IX). Under the uniform and bond-aware models this changes nothing, because both charge a site by the number of qubits its bonds occupy, $\lceil \log_2 \chi \rceil$, so $\chi = 12$ costs what $\chi = 16$ does and reaches less. The question is then whether a compiler could charge intermediate bonds less. The most generous answer is the isometry lower bound evaluated at the actual dimensions, $\lceil (2ab - a^2 - 3\log_2 b - 1)/4 \rceil$ CNOTs for an isometry from dimension $a = D_{k+1}$ to $b = 2D_k$, which equals the bound above at powers of two. With it, and every $\chi \leq 64$ available to the MPS alone, the chain wins fall from 8 to 0 of 63 and the 4×3 wins from 3 to 2, while on 4×4 the measured route still wins 13 of 14; with the transferred step (Sec. IX)

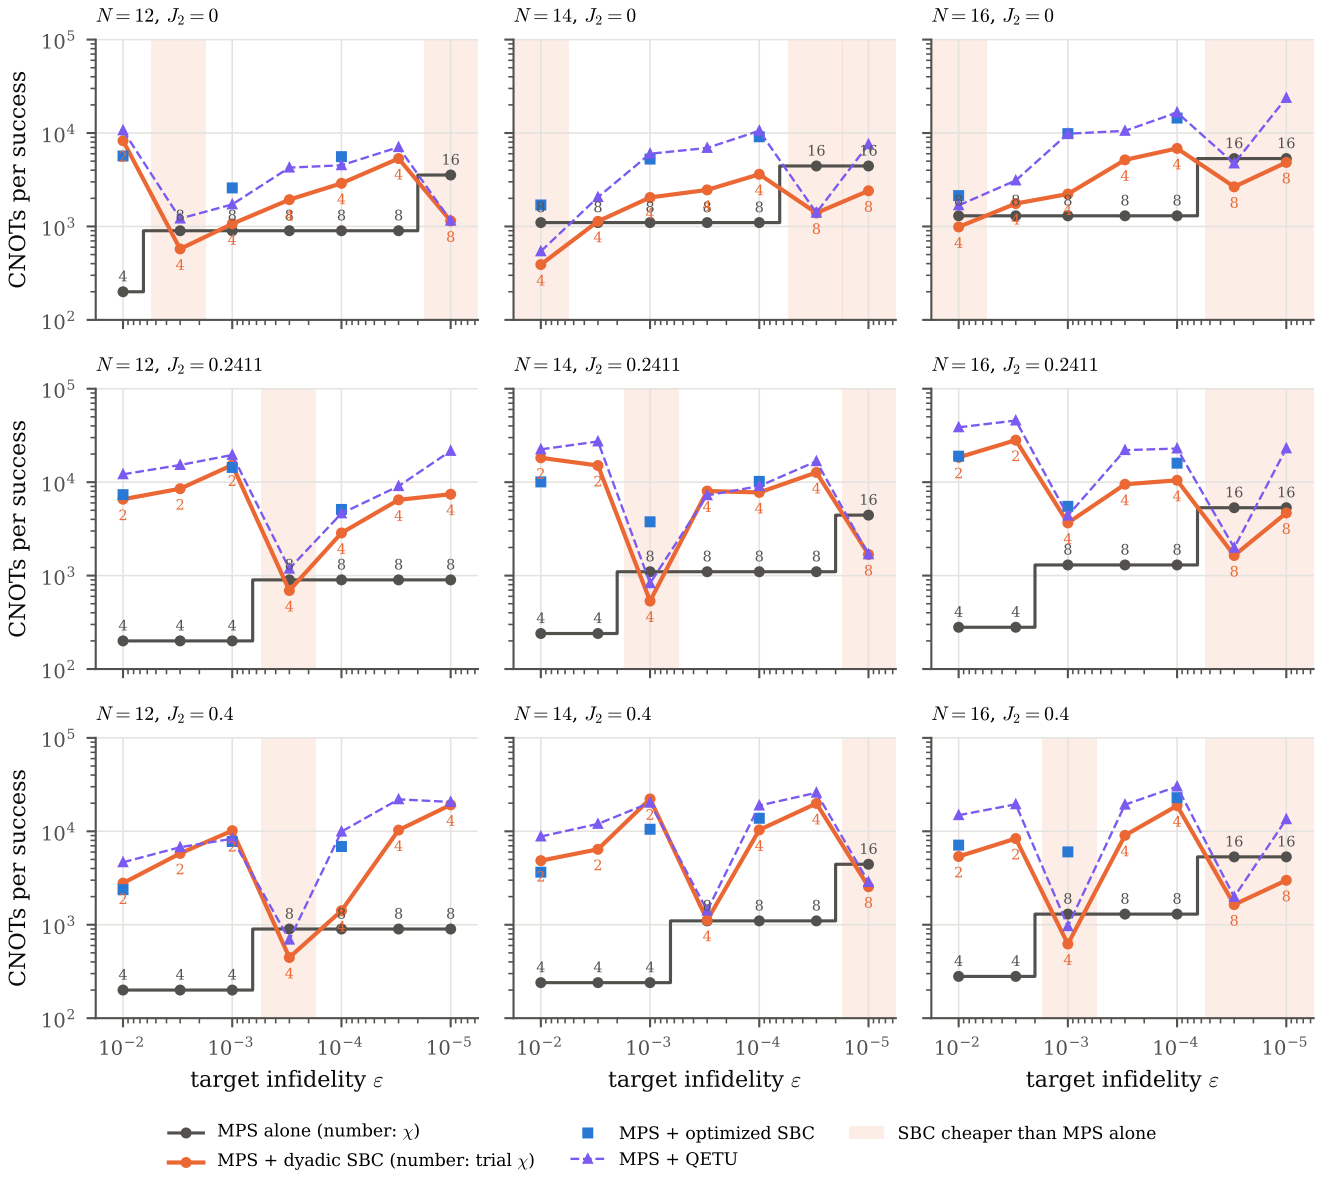


FIG. 10. Measured cheapest route to infidelity ϵ on chains of $N = 12, 14, 16$ (columns) at $J_2 = 0, 0.2411, 0.4$ (rows), all by exact simulation. Grey: cheapest MPS that meets $1 - \gamma \le \epsilon$ alone (number: χ). Orange: cheapest smaller MPS followed by the dyadic SBC filter (number: trial χ); blue squares: optimized SBC (designed at $\epsilon = 10^{-2}, 10^{-3}, 10^{-4}$ only); purple: QETU. Shaded: the SBC route is cheaper. Cost is (trial + filter CNOTs)/p. Sources: `results/emp_actual1d.jsonl` and, for $N = 16$, $J_2 = 0.2411$, `results/emp_actual.json`.

it wins 1 of 14 on the $\chi = 2^k$ grid and 0 with every χ. The measured 4×4 advantage is therefore robust to the MPS cost model and bond-dimension grid, and the implementable one is not: it holds with constructive MPS costs and disappears against a lower bound. *A third cylinder.* The 4×4 result rests on one system size, so we repeated the implementable comparison on 4×5 (20 qubits; `run_cross20.py`) for the same 14 targets: the variational MPS alone at every $\chi \le 256$ against the exact ground state, and the planned filter on the trial χ with the cheapest certified plan, simulated at the step transferred from 4×4. The measured crossover, which needs the full spectrum, is out of reach at this size. The transferred step met the target in all 28 circuits (and in 28 without halving). The filtered route wins 12 of 14 targets under the uniform model and 11 under the bond-aware one, against 12 and 5 on 4×4; against the dimension lower bound it wins 6 on the $\chi = 2^k$ grid, 4 with every $\chi \le 64$ and 2 with every $\chi \le 256$. The implementable advantage therefore grows from 4×4 to 4×5 under constructive costs, and against the lower bound it is small but no longer zero.

Other trials. All trials here are snake MPS, a deliberately simple and, in two dimensions, weak choice; the statements above compare filtering a snake MPS with compiling a better snake MPS. MPS with explicit $U(1)$ or $SU(2)$ symmetry, whose circuits act on smaller blocks, two-dimensional ansätze such as PEPS, and variational circuits were not tried. A better trial would help both routes, and it would move the comparison toward the tight targets, where the filter's $\ln(1/\epsilon)$ depth is its advantage.

These are small cylinders, at sizes where exact simulation is possible. With certified plans in place of measured ones, the same comparison extended to 4×12 finds the filtered route to be the only MPS-based route at $\epsilon = 10^{-4}$ on the larger cylinders, where generic preparation is out of reach. For the two-dimensional J_1–J_2 model, resource estimates for phase estimation have assumed the state preparation to be negligible [29]; the 4×4 results suggest that a compiled low-χ MPS plus a short filter could meet that assumption, provided its Trotter step can be chosen without the worst-case bound (Sec. IX).

VII. COMPARISON WITH A CHEBYSHEV FILTER (QETU)

To compare with a polynomial filter under the same accounting, the pipeline also designs the quantum eigenvalue transformation of unitaries (QETU) of Dong, Lin and Tong [9]. It uses the same primitive, d queries of $e^{-i\tau H_s \otimes Z_a}$, the controlled-reversal form also used by the rodeo algorithm [10], but interleaves them with ancilla rotations and post-selects only once. On every eigenvector the ancilla then undergoes quantum signal processing in $x = \cos(\tau E)$, so the ancilla-0 block is $p(\cos \tau H_s)$ for any real even polynomial p of degree d with $|p| \le 1$ on $[-1, 1]$ [9]. With $\tau = \pi/2$ the scaled spectrum maps to $x \in [0, 1]$ and the stop band to $[0, x_s]$, $x_s = \cos(\pi\Delta/2)$. We take the Chebyshev filter

$$p(x) = \frac{T_d(x/x_s)}{T_d(1/x_s)}, \qquad \eta = \frac{1}{T_d(\cos(\tau w_0)/x_s)}, \tag{15}$$

which is extremal: no polynomial of degree d bounded by 1 is smaller on the stop band relative to its value at the pass band. It is certified in closed form (p is monotone on $[x_s, 1]$ and equioscillates on $[-x_s, x_s]$), has $p(0) = 1$ so $p_g = \gamma F_{\text{pass}}^2 \approx \gamma$, and needs $T = d\tau \approx \Delta^{-1} \ln(2/\eta)$, two to three times less than the dyadic ladder. The Trotter bound, step allocation and Theorem 1 carry over unchanged, with one pulse per query and two ancilla rotations per query (the phase and the identity term). The QSP phases are computed classically; our cost model does not include that step. For second-order formulas $S_2(-\tau) = S_2(\tau)^\dagger$, so the Trotterized circuit implements exactly $p(\tilde{X})$ with $\tilde{X} = \frac{1}{2}(S_2(\tau) + S_2(-\tau))$, which is how we simulate it; the test suite checks this against an explicit ancilla circuit with QSP phases.

Certified. The Chebyshev QETU filter, planned with the same Trotter, synthesis and cost model, is feasible in all 36 second-order cases (0 failures) and is a median 3.3× (range 2.1–3.9) cheaper than the dyadic schedule. Against the optimized SBC design, where that design exists, the ratio is a median 1.00 (range 0.85–4.3, 29 cases): the two cost the same to within the spread, because both reach $T \approx \Delta^{-1} \ln(1/\eta)$ up to a constant and the Trotter cost depends only on T. At $N = 48$, $J_2 = 0.2411$ QETU needs 8.2×10^7 CNOTs per success against 2.7×10^8 for the dyadic plan (Fig. 7). Simulated as in Sec. III B, the certificate held in 24 of 24 QETU plans ($N \le 14$, both targets), with a median looseness of 7.5×; its leakage term is tight (median 1.3×) because the Chebyshev filter equioscillates across the stop band. In the controlled study of Sec. IV, the certified QETU plan is a median 0.30 times the certified dyadic one and cheaper in all 222 cases, and it is the cheapest certified filter in 36, 49 and 74 of 74 cases at $\epsilon = 10^{-2}$, 10^{-3}, 10^{-4} (the optimized SBC design takes the rest). Certified, the SBC filter therefore offers no cost advantage over a polynomial filter; what it offers is the absence of a phase-finding step and a measurement after every pulse, which an early-abort protocol (not modeled here) could exploit.

Measured. The gap/overlap study simulates both filters on every case (Figs. 4–6, bottom rows and purple points). Measured, the dyadic SBC filter is the cheaper one at loose targets: QETU is cheaper in 0 of 74 cases at $\epsilon = 10^{-2}$, 9 of 74 at 10^{-3}, but 56 of 74 at 10^{-4}; over all cases QETU costs a median 1.26× the dyadic filter. The reason is the Trotter step, not the filter. QETU needs less evolution time (a median 0.65× the dyadic time), but it must take one fine Strang step per short query, $\delta t = \tau/W \propto 1/N$, so its step count is a median 1.40× the dyadic one and its CNOTs grow as $N^{2.22}$ against $N^{1.30}$. The long dyadic pulses tolerate steps of order $1/J_1$, but these shrink as roughly $\epsilon^{0.3}$, and by $\epsilon = 10^{-4}$ the shorter QETU filter wins. QETU's measured cost is also the more regular of the two (Table IV): its fit holds to 1.04×, because its degree grows by two at a time where the dyadic ladder grows in whole pulses.

Query length. The measured comparison above uses $\tau = \pi/2$, the choice of Dong, Lin and Tong [9], and since one Strang step per query always sufficed, τ fixes QETU's Trotter step. It is not forced: any $\tau \le \pi/(1+\Delta)$ keeps every excited level inside $|\cos \tau E| \le x_s = \cos \tau\Delta$, and the even Chebyshev filter applies unchanged with fewer, longer queries. We therefore redid every measured QETU search (467 cases: the controlled study and the crossovers) at $\tau = 3\pi/4$ and at the largest admissible $\tau = \pi/(1 + \Delta)$, which is also the dyadic base time (`run_qetu_tau.py`). The total evolution time barely changes, but the query count falls and one Strang step per query still suffices, so the measured cost falls to a median 0.68 (range 0.29–4.28) of the $\tau = \pi/2$ cost at $\tau = \pi/(1 + \Delta)$, and 0.75 at $3\pi/4$. That is enough to change the measured ranking (Table X): with QETU at its best τ the SBC filters keep $\epsilon = 10^{-2}$ (19 dyadic and 48

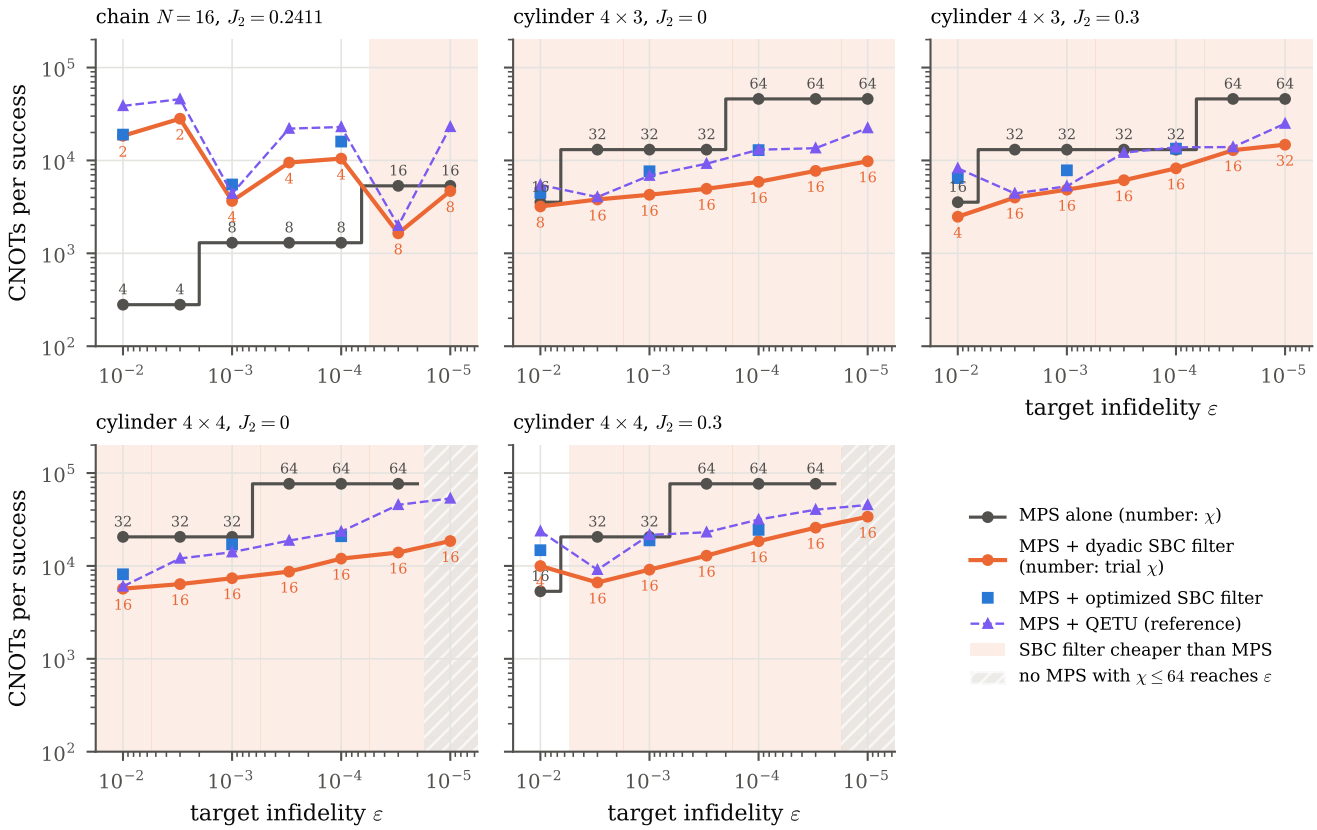


FIG. 11. Measured cheapest route to infidelity ϵ on the 16-site chain and on small cylinders. Grey: the cheapest MPS that meets the target alone (number: χ). Orange: the cheapest smaller MPS followed by the dyadic SBC filter (number: trial χ); blue squares: optimized SBC; purple: QETU. Shaded: the filtered route is cheaper; hatched: no MPS with $\chi \leq 64$ reaches ϵ. Every cost is (trial + filter CNOTs)/p from exact simulation. Source: `results/emp_actual.json`.

TABLE VIII. Crossovers under three MPS cost models, applied to the trial in the filtered route and to the MPS alone. Uniform: the staircase bound used in the study, $(N-b)$ quantum-Shannon unitaries on $b+1$ qubits. Bond-aware: each site unitary sized by its actual bonds, with exact generic state preparation [38] as an extra baseline. Lower bounds: the isometry bound [37] per site and for generic state preparation. Entries: targets at which MPS + SBC filter is the cheapest route (largest advantage in parentheses), measured and certified. Sources: `results/emp_actual*.json*`, `results/cert_crossover.jsonl`.

	9 chains		4×3		4×4	
MPS cost model	measured	certified	measured	certified	measured	certified
uniform QSD bound (study)	18/63 (3.2×)	0/63	14/14 (7.8×)	0/14	13/14 (8.8×)	2/14
bond-aware QSD, or generic prep.	16/63 (2.9×)	0/63	5/14 (1.4×)	0/14	13/14 (5.7×)	0/14
lower bounds (isometry, generic state)	8/63 (2.1×)	0/63	3/14 (1.7×)	0/14	13/14 (3.2×)	0/14

optimized against 7 QETU of 74), the two families split $\epsilon = 10^{-3}$ (QETU cheapest in 38 of 74 cases), and QETU is the cheapest filter in 65 of 74 at 10^{-4}. The statement that SBC wins measured at $\epsilon \geq 10^{-3}$ therefore holds only for $\tau = \pi/2$; the fair statement is that SBC wins at loose targets, $\epsilon \approx 10^{-2}$, the two are even at 10^{-3}, and the Chebyshev filter wins from 10^{-4} on. In the crossovers, QETU at its best τ beats the MPS alone at 14 of 63 chain and 26 of 28 cylinder targets but is the cheapest route at only 2 chain and 1 cylinder targets, so the crossover results of Secs. V D and VI change at three targets of 91 when QETU is given its best τ. Two remaining differences in accounting favor neither side much: we cost QETU in the controlled form, because no Pauli string anticommutes with a Heisenberg bond (Sec. II), and we leave out the classical phase-finding step, which costs no quantum gates.

Appendix C analyzes what sets the Chebyshev filter's

TABLE IX. Intermediate bond dimensions for the MPS alone. Targets at which MPS + filter is the cheapest route, measured (minimum over trial χ, design and step) and with the planned filter and transferred step, when the MPS alone may use only $\chi = 2^k$ or every $\chi \le 64$ (adding 3, 5, 6, 10, 12, 20, 24, 40, 48). Dimension lower bound: the isometry bound at the actual bond dimensions, applied to the MPS alone, and at powers of two to the trial. The transfer rule has no smaller system for 4×3. Sources: `results/chi_mid.json`, `results/emp_actual*.json*`, `results/aposteriori.jsonl`.

	9 chains (63)		4×3 (14)		4×4 (14)	
MPS cost model, χ available to the MPS alone	measured	transfer	measured	transfer	measured	transfer
uniform QSD bound, $\chi = 2^k$	18	0	14	–	13	12
uniform QSD bound, all $\chi \le 64$	18	0	14	–	13	12
bond-aware QSD, $\chi = 2^k$	16	0	5	–	13	5
bond-aware QSD, all $\chi \le 64$	16	0	5	–	13	5
dimension lower bound, $\chi = 2^k$	8	0	3	–	13	1
dimension lower bound, all $\chi \le 64$	0	0	2	–	13	0

TABLE X. Cheapest measured filter in the controlled study (74 cases per target; dyadic/optimized/QETU) with QETU at $\tau = \pi/2$ and at its best query length $\tau \in \{\pi/2, 3\pi/4, \pi/(1+\Delta)\}$. Source: `results/qetu_tau.jsonl`.

target	$\tau = \pi/2$	best τ
$\epsilon = 1 \times 10^{-2}$	23/51/0	19/48/7
$\epsilon = 1 \times 10^{-3}$	16/57/1	4/32/38
$\epsilon = 1 \times 10^{-4}$	5/32/37	0/9/65

degree, by simulation for $N = 8$–16.

VIII. MEASURED VERSUS CERTIFIED

Every comparison above has two versions, and they do not always agree (Table XI). The measured cost is not implementable as it stands: it picks the shortest filter and the coarsest Trotter step that reach ϵ for the given trial, which needs the exact ground state. It is the cost an ideal a-posteriori procedure could approach, and it shows what the circuit actually has to do. The certified cost is what the pipeline plans from Theorem 1 without that knowledge, and can be run as is. The two differ mainly through the Trotter step (Sec. IV): the certified step is 60–75 times finer, and its overhead grows with the filter length.

This difference decides which filter is cheapest. Measured, the long SBC pulses tolerate steps of order $1/J_1$ and win at loose targets; certified, every filter pays the same worst-case step per unit time, the shorter Chebyshev filter wins more often, and the optimized SBC design is the only one that competes. It also decides the end-to-end comparison. Certified, the filtered route is 1–2 orders of magnitude more expensive than measured. Against compiling a better MPS it loses at all 63 chain targets (median 152.3× the MPS cost, never below 13.5×) and at every cylinder target where some $\chi \le 64$ reaches ϵ (median 5.2×), with one exception: QETU on 4×3, $J_2 = 0$, at $\epsilon = 10^{-4}$, which is 1.2× cheaper and where the SBC filter ties; with the stronger baselines of Table VIII it wins none. It remains the only MPS-based route where no MPS reaches the target, but not cheaper than exact generic preparation there. The measured cylinder advantage is therefore a statement about what the circuit needs, which an implementation can claim only with a tighter error bound, for example one restricted to the low-energy subspace [19, 31], or with an a-posteriori choice of the Trotter step, which we test in Sec. IX. Two smaller oracle elements are the exact gap and overlap used to design the measured filters; the sensitivity study of Sec. V B shows that an overestimated gap is the dangerous error, so a margin on E_1 is needed in practice.

IX. CHOOSING THE TROTTER STEP A POSTERIORI

The measured costs use two oracle elements, the filter length and the Trotter step, and the certified step is 60–75 times finer than needed. We test whether the step can be chosen without the bound and without the ground state, keeping the filter as planned by Theorem 1 from the estimated gap and overlap. For every chain and cylinder target of the crossovers and both schedules (SBC "auto" and QETU), the trial χ whose certified plan is cheapest is filtered with its planned circuit, Trotterized on a common grid of physical steps $\delta t = 4 \cdot 2^{-n}/J_1$ ($k_i = \lceil t_i/\delta t \rceil$ Strang steps per pulse or query), 182 cases in all (`run_aposteriori.py`). At each step we record the infidelity (the oracle check), the success probability, the post-selected energy, and the variances of the XX, YY and ZZ parts of H, from which the shot cost of measuring that energy follows. With the oracle step, the planned filter costs a median 0.09 (SBC) and 0.12 (QETU) of its certified cost: the step, not the filter, is almost the whole

TABLE XI. Each comparison measured (shortest filter and coarsest Trotter step that reach ϵ in exact simulation; uses the ground state) and certified (planned from Theorem 1; implementable). Controlled study: 74 cases per target. Crossovers: targets 10^{-2}–10^{-5}; certified SBC uses the pipeline's automatic choice of dyadic or optimized, QETU in parentheses. Sources: `results/gapov*.jsonl`, `results/emp_actual*.json*`, `results/cert_crossover.jsonl`.

comparison	measured	certified
Cheapest filter, $\epsilon = 1 \times 10^{-2}$ (dy/opt/QETU)	23/51/0	0/38/36
Cheapest filter, $\epsilon = 1 \times 10^{-3}$ (dy/opt/QETU)	16/57/1	0/25/49
Cheapest filter, $\epsilon = 1 \times 10^{-4}$ (dy/opt/QETU)	5/32/37	0/0/74
MPS + filter beats MPS alone, 9 chains	18 of 63	0 of 63 (QETU 0)
MPS + filter beats MPS alone, 4 cylinders	27 of 28	2 of 28 (QETU 3)

certified overhead. Two rules that a user could apply are then evaluated (Table XII).

Energy convergence. Halve δt until the post-selected energy changes by at most $\epsilon\Delta/2$ between successive steps, and use the finer one. Since the leaked weight sits at least Δ above E_0, an infidelity ϵ moves the energy by at least $\epsilon\Delta$, so the rule is conservative by design. It met the target in 182 of 182 cases, at a median 0.22 (SBC) and 0.23 (QETU) of the certified cost, 1.9 and 1.9 times the oracle step. Its price is the calibration: resolving $\epsilon\Delta/2$ with the three measurement bases takes a median 2×10^{10} post-selected shots per case (about 10^7 at $\epsilon = 10^{-2}$ and 10^9 at 10^{-3}, growing as ϵ^{-2} to $\sim 10^{13}$ at 10^{-5}), which pays off only if the same circuit is then used for $\gtrsim 6\times10^9$ preparations. It is a reliable rule whose cost is in shots rather than gates, and at tight targets it is not a practical alternative to the transferred step: from $\epsilon = 10^{-3}$ on, the calibration needs $\gtrsim 10^9$ post-selected shots before the circuit is used once.

Size transfer. Use half the coarsest passing step of the next-smaller system with the same J_2, ϵ and schedule ($N = 12 \to 14 \to 16$, $4 \times 3 \to 4 \times 4$), which needs exact simulation only at the smaller size. The rule met the target in 112 of 112 cases, at a median 0.13 (SBC) and 0.18 (QETU) of the certified cost, within a factor of about two of the oracle step.

Why the step should transfer, and how far. A Strang step is $e^{-i\,\delta t\,H_{\rm eff}}$ with $H_{\rm eff} = H + \delta t^2 V + O(\delta t^4)$, V a sum of local nested commutators, so the Trotterized filter is (up to the rounding of k_i) the exact filter applied in the eigenbasis of $H_{\rm eff}$. Two errors follow, and both are extensive. The ground state of $H_{\rm eff}$ differs from that of H by an infidelity $\sim \delta t^4 \sum_n |\langle n|V|0\rangle|^2/(E_n - E_0)^2 \propto N\delta t^4$, so a fixed infidelity needs $\delta t \propto N^{-1/4}$; and the ground energy moves by $\delta t^2\langle 0|V|0\rangle \propto N\delta t^2$ relative to the window the filter was planned for, so a fixed displacement needs $\delta t \propto N^{-1/2}$. The step is therefore *not* independent of N; what the rule relies on is that the step shrinks slowly, so that the factor of two covers a size ratio of up to 16 under the first mechanism and 4 under the second. In the simulated range the coarsest passing step fell by a factor of two in 14 of the 112 size steps and never rose by more. Transferring the step between values of J_2 at the same size, which changes the local terms and their commutators (and the gap by at most 16%), is a test where the step is known to move: the coarsest passing step differs in 90 of the 308 pairs, and the halved step met the target in 308 of 308. Gaps far from the calibration system's, such as the 0.07 of the natural chains at $N = 48$, were not tested.

A size beyond the calibration. To test the rule one step further than any earlier simulation, we simulated the planned circuits on a 20-site chain (step from $N = 16$) and on a 4×5 cylinder (step from 4×4), for every J_2, $\epsilon = 10^{-2}$, 10^{-3}, 10^{-4} and both schedules, with DMRG inputs, the trial χ used at the smaller size and the exact ground state (30 circuits, statevector simulation of 20 qubits; `run_transfer20.py`, with each bond's three commuting rotations applied as one exact SWAP power, which agrees with the term-by-term product formula to 10^{-14}). The transferred step met the target in 30 of 30 cases at a median 0.12 of the certified cost (Table XIII), and in 28 of 30 the coarsest passing step of the smaller system, without the factor of two, already sufficed. This validates one size step of 25%; it does not validate the factor of three from 4×4 to 4×12.

End to end. With the planned filter and the transferred step, neither element of the route uses the target system's ground state, and on 4×4 it still beats the MPS alone at 12 of 14 targets (up to 2.6×) with the uniform MPS cost model and 5 of 14 (up to 1.8×) with the bond-aware one, against none certified; with the energy rule the counts are 7 and 3. On chains it wins none of the 63 targets with either rule, and on 4×3 only with the energy rule and the uniform model (5 of 14).

Projections for larger cylinders. Applied to the planned filters of 4×6, 4×8 and 4×12 (DMRG inputs, step transferred from 4×4; Table XIV), the same accounting favors the filtered route at 15 of 18 targets with the bond-aware MPS cost (17 with the uniform one), and it is the only MPS-based route at 4 of them. These are projections, not results: the step is transferred over a size ratio of 1.5–3 that we could not simulate, the success probability is taken as $p_g = \gamma F_{\rm pass}^2$, and the overlaps come from DMRG. To bracket the step, the table also gives the advantage with the unhalved 4×4 step and with the transferred step shrunk by a further $\sqrt{16/N}$ (the energy-shift scaling above): with the second the count falls to 13 of 18, and 5 of the finite advantages change sign within the bracket. Only the advantages above about 1.5 at $\epsilon = 10^{-3}$ and the four "only" entries survive the whole

range.

X. DISCUSSION

What makes the filter advantageous. The controlled study separates three regimes. (i) The gap determines the coherent circuit: evolution time $\propto L/\Delta$, independent of N at fixed gap, and CNOTs and R_z per shot that scale as a power of $1/\Delta$. A system with a large gap in absolute units is the natural target, and the filter's depth is what hardware with a coherence or logical-error budget has to accommodate. (ii) The trial overlap enters the circuit only through a logarithm and the expected cost through $1/\gamma$: the filter is tolerant of a poor trial, paying in repetitions rather than in depth. This is where it has an edge over compiling a better trial, whose cost grows with the bond dimension needed to reach ϵ directly. (iii) The target enters through L and, for SBC, through the Trotter step, $\delta t \propto \epsilon^{0.3}$; tight targets are where a compiled MPS runs out first. For the chain, the MPS infidelity falls quickly with χ, so the filter pays off only as a last-mile polisher, on a trial within a factor of about 8.8 of the target (Sec. V D); on the 4×4 cylinder, where a snake MPS struggles, it pays off at 13 of 14 targets against the stronger baselines, while on 4×3 exact generic preparation is already cheap enough to compete (Sec. VI).

Dyadic, optimized or Chebyshev. Measured, the SBC schedules beat the Chebyshev filter at $\epsilon \approx 10^{-2}$, because their long smooth pulses tolerate Trotter steps of order $1/J_1$; with its longest admissible query, $\tau = \pi/(1+\Delta)$, the shorter Chebyshev filter ties at 10^{-3} and wins from 10^{-4} on (Table X), and certified it is usually the cheaper one (Sec. VIII). Among the SBC schedules, the optimized design is the cheaper at $\epsilon \leq 10^{-3}$ in most cases, measured (median 0.89× dyadic) and even more so certified (0.35×), but it needs a numerical design that occasionally fails and must be repeated; the dyadic ladder is the robust default and the natural fallback, which is how the pipeline's "auto" mode uses them.

Amplitude amplification. All costs here are repeat-until-success, C/p per success. Amplitude amplification [35] would instead run the filter coherently about $\pi/(4\sqrt{p})$ times, each round costing the trial and filter circuits and their inverses, for an expected cost near $(\pi/2)C/\sqrt{p}$: a saving of about $(\pi/2)\sqrt{p}$, i.e. worthwhile only for $p \lesssim 0.4$. At the lowest overlap of the controlled study ($p \approx \gamma = 0.03$) it would cut the cost per success by about 3.7×, at the price of a coherent circuit about six times deeper, which works against the depth budget the gap already sets. Two further costs apply to the cosine filter. The SBC circuit measures the ancilla after every pulse; a coherent version needs one ancilla per pulse (or the pulses combined into a single post-selection, as in QETU), and with p known only as an estimate a fixed-point variant [36] is needed, which costs a further logarithmic factor. Where the filter wins in the crossovers, the trial is good and $p \geq 0.77$, so amplification would not help; it matters in the low-overlap regime, where it changes the $1/\gamma$ of the cost laws into $1/\sqrt{\gamma}$. An early-abort protocol that stops at the first failed pulse is the cheaper complement for SBC and is likewise not modeled.

What is guaranteed. Theorem 1 holds whenever its inputs are right. Of these, the spectral upper bound (Weyl) and the filter certificate are rigorous, and the Trotter constants are exact up to the Lanczos tolerance of their largest local norms. The interval for E_0, the lower bound on E_1 and the overlap γ come from DMRG and are well-converged estimates, not bounds; for $N \leq 14$ ED confirms them, beyond that the certificate is conditional on all three. The gap is the critical input: an overestimate of 10–25% already broke the target for the optimized and Chebyshev filters (Table VI). Rigorous lower bounds on E_1 (e.g. Temple-type bounds) would remove the gap condition, and with them the Eckart inequality $1 - \gamma \leq (\langle H \rangle_\psi - E_0)/(E_1 - E_0)$ would remove the overlap condition. "Certified" in this paper therefore means *planned from Theorem 1*: the guarantee is unconditional only for $N \leq 14$ with ED inputs (and up to the Lanczos tolerance of the Trotter constants), and for larger systems it is conditional on three unverified inputs. Because the gap bias of DMRG runs in the dangerous direction and the optimized and Chebyshev plans have no slack, a deployable plan should carry an explicit margin δE_1; in the instance of Sec. V B a margin that assumes 0.5 of the gap costs 2.9×, small beside the Trotter overhead.

Where the certified cost goes. The certified filter length is within a factor of about two of what simulation needs; the certified Trotter step is 60–75 times finer than needed, and the certified step count grows as L^2 rather than L. Tighter Trotter bounds, for example restricted to the low-energy subspace [19, 31], would recover most of the 62-fold median gap between certified and measured cost; a better filter design would not. The a-posteriori rules of Sec. IX already recover most of it in practice: transferring the step from a smaller, classically simulable system costs about twice the oracle step and met the target in every case we simulated, across J_2 and for one 25% size step to 20 qubits. That record covers size ratios of at most 1.33 and is not evidence for the 4×6–4×12 projections: the step shrinks with N (as $N^{-1/4}$ to $N^{-1/2}$ by the argument of Sec. IX), so the factor of two runs out at a size ratio between 4 and 16, and the projections rest on the bracket of Table XIV, not on the simulations; checking energy convergence on the device is rigorous in spirit but costs ϵ^{-2} shots.

Limitations. Noiseless simulation; repeat-until-success rather than early abort; all-to-all connectivity (a nearest-neighbor layout would add SWAPs, mostly for the ancilla, which couples to every bond); trial circuits costed by a generic decomposition bound rather than compiled (Table VIII brackets this); exact simulation and the controlled study only up to $N = 14$–16 and 16-site cylinders, sizes at which classical methods are exact, so the fitted exponents are descriptive and every claim about larger systems rests on certified plans or on the step transfer of

TABLE XII. Planned (Theorem 1) filters with an a-posteriori Trotter step. Columns: cases; median cost with the oracle step over the certified cost; energy-convergence rule (targets met, median cost over certified, median post-selected shots for the calibration); size-transfer rule (targets met, median cost over certified). Source: `results/aposteriori.jsonl`.

		oracle	energy convergence			size transfer	
schedule, system	cases	cost/cert.	met	cost/cert.	shots	met	cost/cert.
SBC, chain, $N = 12$	21	0.08	21/21	0.22	2×10^{10}	–	–
SBC, chain, $N = 14$	21	0.07	21/21	0.22	9×10^{10}	21/21	0.13
SBC, chain, $N = 16$	21	0.08	21/21	0.19	3×10^{11}	21/21	0.12
SBC, 4×3	14	0.19	14/14	0.32	1×10^{11}	–	–
SBC, 4×4	14	0.14	14/14	0.23	1×10^{11}	14/14	0.16
QETU, chain, $N = 12$	21	0.10	21/21	0.22	3×10^{9}	–	–
QETU, chain, $N = 14$	21	0.10	21/21	0.22	4×10^{9}	21/21	0.17
QETU, chain, $N = 16$	21	0.10	21/21	0.21	6×10^{9}	21/21	0.18
QETU, 4×3	14	0.21	14/14	0.34	6×10^{9}	–	–
QETU, 4×4	14	0.15	14/14	0.24	1×10^{10}	14/14	0.21

TABLE XIII. Size transfer beyond the calibration sizes, simulated: planned filter at 20 qubits with the step of the smaller system. Columns: cases; target met with the smaller system's coarsest passing step (no safety factor); target met with the transferred (halved) step; largest infidelity over ϵ at the transferred step; median cost over certified. Source: `results/transfer20.jsonl`.

	cases	unhalved	transferred	max ℓ/ϵ	cost/cert.
SBC, chain, $N = 20$ (from 16)	9	7/9	9/9	0.26	0.09
QETU, chain, $N = 20$ (from 16)	9	9/9	9/9	0.41	0.11
SBC, 4×5 (from 4×4)	6	6/6	6/6	0.29	0.11
QETU, 4×5 (from 4×4)	6	6/6	6/6	0.28	0.22

Sec. IX; the gap range that can be held fixed across N is narrow (0.30–0.55, set by the finite-size gap at $N = 8$–14) and lies in the easy regime, well above the gaps near 0.07 of the natural chains at $N = 48$, which enter only through the certified plans of Sec. V; the measured costs are minima over filter designs, Trotter steps and (in the controlled study) candidate ladders, a selection that the MPS-alone side matches only by choosing the cheapest sufficient χ (Table IX quantifies how much of the advantage this selection carries); E_1 is taken as the lowest $S^z = 1$ state, which is the first excitation in every ED case here but an assumption for larger N.

CODE AND DATA AVAILABILITY

The `sbpipe` package, the study scripts, all raw results (`sbpipe_study/results/*.json`) and the sources of this paper are available at https://github.com/aaronseymour7/FilterPipeline.

Appendix A: Proof of Theorem 1

Write the trial state as $|\psi\rangle = \sqrt{\gamma'}|g\rangle + \sqrt{1-\gamma'}|\chi\rangle$ with $|g\rangle$ in the ground space (energy $e_0 \in [0, w_0]$ in scaled units), $|\chi\rangle$ orthogonal to it and $\gamma' \geq \gamma$.

Lemma 1 (leakage)**.** *For the exact filter, $v = F(H_s)|\psi\rangle$ has weight outside the ground space at most ℓ, and $\|v\|^2 \geq p_g$.*

Proof. $F(H_s)$ is a function of H_s, so $v = \sqrt{\gamma'}F(e_0)|g\rangle + \sqrt{1-\gamma'}F(H_s)|\chi\rangle$ with the two parts orthogonal. Every component of $|\chi\rangle$ has energy in $[\Delta, 1]$, so $\|F(H_s)\chi\| \leq \sup_{[\Delta,1]}|F| \leq \eta F_{\text{pass}} \leq \eta|F(e_0)|$. The excited weight is $x/(a+x)$ with $a = \gamma' F(e_0)^2$, $x = (1-\gamma')\|F\chi\|^2 \leq (1-\gamma')\eta^2 F(e_0)^2$; it is increasing in x and, after dividing by $F(e_0)^2$, decreasing in γ', which gives ℓ. Finally $\|v\|^2 \geq a \geq \gamma F_{\text{pass}}^2$. □

For pure states the trace distance is $\sin\angle$, so $D(\hat{v}, |g\rangle) = \sqrt{\ell'} \leq \sqrt{\ell}$ ($\hat{v} = v/\|v\|$).

Lemma 2 (implementation error)**.** *Let $B_i = \frac{1}{2}(e^{-i\phi_i}\tilde{S}_i(t_i) + e^{i\phi_i}\tilde{S}_i(-t_i))$ be the ancilla-0 block of the implemented pulse and $C_i = \cos(H_s t_i + \phi_i)$. Then $\|\prod_i B_i - \prod_i C_i\| \leq \epsilon_U$.*

TABLE XIV. Larger cylinders, projected: planned filter with the Trotter step transferred from 4×4, not simulated. Expected CNOTs per success of the cheapest filtered route and its advantage over the MPS alone (bond-aware cost; "only": no $\chi \leq 64$ MPS reaches ϵ). In parentheses: the advantage with the step shrunk by a further $\sqrt{16/N}$ and with the unhalved 4×4 step; the range covers the step only, not the DMRG overlaps or p_g. Sources: `results/ext_regime_cyl*.json`, `results/aposteriori.jsonl`.

	$\epsilon = 10^{-2}$		10^{-3}		10^{-4}	
cylinder	CNOT	adv.	CNOT	adv.	CNOT	adv.
4×6, $J_2 = 0$	2.2×10^4	1.32 (1.21–1.56)	4.6×10^4	2.28 (1.97–3.50)	8.6×10^4	1.22 (0.88–1.82)
4×6, $J_2 = 0.3$	2.4×10^4	1.22 (1.11–1.51)	5.8×10^4	1.81 (1.71–2.54)	1.1×10^5	0.92 (0.72–1.49)
4×8, $J_2 = 0$	4.0×10^4	1.09 (0.96–1.32)	7.1×10^4	2.32 (1.83–3.46)	1.7×10^5	only
4×8, $J_2 = 0.3$	4.6×10^4	0.96 (0.85–1.21)	1.0×10^5	1.65 (1.35–2.62)	2.1×10^5	only
4×12, $J_2 = 0$	7.9×10^4	3.66 (2.76–4.62)	1.9×10^5	1.51 (1.22–1.80)	4.0×10^5	only
4×12, $J_2 = 0.3$	9.0×10^4	0.82 (0.62–1.03)	1.7×10^5	1.65 (1.08–2.48)	3.4×10^5	only

Proof. $C_i = \frac{1}{2}(e^{-i\phi_i}e^{-it_iH_s} + e^{i\phi_i}e^{it_iH_s})$, so $\|B_i - C_i\| \leq \frac{1}{2}(\|\tilde{S}_i(t_i) - e^{-it_iH_s}\| + \|\tilde{S}_i(-t_i) - e^{it_iH_s}\|) \leq \epsilon_i$, where ϵ_i is the product-formula bound (symmetric in the sign of t) plus the synthesis error of the pulse's rotations. Since $\|B_i\|, \|C_i\| \leq 1$, telescoping gives $\|\prod B_i - \prod C_i\| \leq \sum_i \epsilon_i$. A synthesized trial circuit prepares $|\tilde{\psi}\rangle$ with $\|\tilde{\psi} - \psi\| \leq R_{\text{trial}}\delta$, and $\|\prod B_i\| \leq 1$, so the total error of the unnormalized output is at most ϵ_U. □

With $\tilde{v} = \prod B_i|\tilde{\psi}\rangle$, $\|\tilde{v} - v\| \leq \epsilon_U$. The distance from v to the line through $\tilde{v}$ is at most $\|v - \tilde{v}\|$, so $\sin\angle(v, \tilde{v}) \leq \epsilon_U/\|v\| \leq \epsilon_U/\sqrt{p_g}$. The triangle inequality for the trace distance gives the bound on D. The states $\hat{v}$, $|g\rangle$ and the normalized output are pure, so $D^2 = 1 - |\langle g|\tilde{\psi}_{\text{out}}\rangle|^2$; since $|g\rangle$ lies in the ground space, $\text{Tr}(\Pi_0\tilde{\rho}) \geq |\langle g|\tilde{\psi}_{\text{out}}\rangle|^2$, which gives (8). (For a mixed output, e.g. under noise, a trace-distance bound alone would not bound the ground-space weight by D^2; this case is outside the theorem.) The success probability is $\|\tilde{v}\|^2 \geq (\|v\| - \epsilon_U)^2 \geq (\sqrt{p_g} - \epsilon_U)^2$. For the energy, the normalized state has weight at least $1 - D^2$ on the ground space and at most D^2 on levels $\leq E_{\text{top}}$, so $\langle H\rangle - E_0 \leq (E_{\text{top}} - E_0)D^2$. □

Appendix B: Dyadic versus optimized SBC, and cost maps

The two SBC schedules reach the same physics by different routes. The dyadic ladder has zero phases, so it passes the ground state with $f(E_0) = 1$ and its length grows in whole pulses. The optimized design (times and phases from SBC's recipe with a certified suppression constraint) is shorter, at the price of $f(E_0) < 1$: its measured success probability is a median 0.88 times γ (as low as 0.52). Measured per shot, the optimized filter uses a median 0.58 times the dyadic evolution time and 0.62 times its Trotter steps; its tolerated step is the same, median $\delta t = 1.24$ at $\epsilon = 10^{-2}$ and 0.30 at 10^{-4}. Per success, after the lower success probability is paid, it is still a median 0.89× the dyadic cost (range 0.58–1.55) and the cheaper one in 169 of 222 cases. Certified, the advantage is larger, a median 0.35× (cheaper in 221 of 222), because the optimized filter's leakage bound is nearly tight while the closed-form dyadic bound is loose (Sec. III B), and the shorter filter needs a looser Trotter budget; its certified cost lies 5–185× (median 26) above the measured one, against 62× for dyadic.

Counting the cheapest of the three filters case by case, the optimized SBC filter wins 51 of 74 cases at $\epsilon = 10^{-2}$ (dyadic the other 23), 57 at 10^{-3} (dyadic 16, QETU 1), and 32 at 10^{-4}, where QETU takes 37 and dyadic 5. The cost laws of Table IV make the trade-off explicit: the optimized measured cost grows more steeply with the gap ($\Delta^{-1.36}$ against $\Delta^{-1.15}$), because its design is matched to the gap, and less steeply with L. The price of the optimized schedule is its design: a nonconvex optimization of about 40 s per case that converged to a filter with $f(E_0)^2 \sim 10^{-3}$ in 21 of 286 measured points and had to be repeated with other seeds; the dyadic ladder needs no design and never fails.

Cost maps.

Two maps make the interplay visible. At fixed N and $\epsilon = 10^{-2}$ (Fig. 12), the contours of the dyadic cost per success run nearly parallel to the gap axis for $g \gtrsim 0.40$: there the overlap alone sets the cost, from 2–6 $\times 10^3$ CNOTs at $\gamma = 0.9$ to 1–4 $\times 10^5$ at $\gamma = 0.03$; below $g \approx 0.40$ the gap starts to matter, and $g = 0.30$ costs 1.3–2.7 times $g = 0.40$ at the same overlap. The optimized filter is cheaper over most of the map, and its contours tilt toward the gap axis over the whole range, as its steeper gap exponent implies. Per shot (not shown) the contours of both turn by ninety degrees and the gap alone sets the circuit. At fixed gap and overlap (Fig. 13), the measured cost per success grows 7–21× (dyadic) from $\epsilon = 10^{-2}$ to 10^{-5} and 1.8–4.6× from $N = 8$ to 14; the contours tilt toward the ϵ axis, so over this range the target matters more than the size. The optimized filter's lead is largest at loose targets and at $\gamma = 0.5$, and shrinks

toward $\epsilon = 10^{-5}$.

Appendix C: What the Chebyshev filter needs

To see how the gap and the overlap act on the Chebyshev filter itself, we replace the certificate by simulation for $N = 8$–16 (675 cases). Each trial state is decomposed on the exact eigenbasis (dense diagonalization per S^z block), so the smallest QETU degree $d_{\min}$ that reaches ϵ for that state is found exactly; at that degree we then find the smallest number of Strang steps per query with the Trotterized circuit simulated on the full state vector, allowing a slightly higher degree when it makes one step per query enough. Three families separate the inputs. In the first, $J_2 = 0.2411$ for every N and the gap closes with N. In the second, J_2 is chosen at each N (by bisection on ED) so that the gap is $E_1 - E_0 = 0.40$ exactly ($J_2 = 0.384, 0.510, 0.539, 0.550, 0.555$ for $N = 8$–16). In the third, J_2 is tuned at $N = 10$, 12 and 14 to gaps from 0.30 to 0.55. For the trial states we use the optimal MPS at $\chi = 1$, 2, 4 and synthetic states $\sqrt{\gamma}|E_0\rangle + \sqrt{1-\gamma}|\xi\rangle$ with γ from 0.03 to 0.9, ξ being the normalized excited part of the $\chi = 2$ MPS, so that the overlap changes while the shape of the excited weight does not. $\epsilon \in \{10^{-2}, 10^{-3}, 10^{-4}\}$.

The degree is set by the gap and the target alone. The measured degrees collapse onto

$$d_{\min}\,\Delta \;\simeq\; \frac{2}{\pi}\cosh^{-1}\frac{1}{\eta}, \qquad \eta^2 = \frac{\gamma\epsilon}{1-\gamma}, \tag{C1}$$

for every family, trial and target (Fig. 14b). This is the certified degree with the true gap: the certificate asks for a median 1.09 times $d_{\min}$. Where the excited weight of the trial sits does not help. The Chebyshev filter is equiripple, so every state above the gap is suppressed by the same factor $1/T_d(1/x_s)$, and a trial whose excited weight lies far above the gap needs the same degree as one whose weight sits at the gap. Over the gap scan $d_{\min} \propto \Delta^{-1.00}$.

The certificate's Trotter count is pessimistic, not its filter. One Strang step per query was enough in all 675 cases, against a certified median of 16 steps per query. Commutator bounds are worst-case over all states, but the filter only has to act well on a state supported near the bottom of the spectrum. As a result the certified cost is a median 22 times what the filter needs (range 7–2×10^2), and the real cost follows the degree:

$$\mathrm{cost}_{\mathrm{real}} \;\approx\; \frac{c_{\mathrm{step}}(N)}{\gamma}\,d_{\min} \;\propto\; \frac{c_{\mathrm{step}}(N)\,W}{g\,\gamma}\,\ln\frac{1-\gamma}{\gamma\epsilon}, \tag{C2}$$

with c_{step} the CNOTs of one step ($\propto N$) and $g = E_1 - E_0$. Fitting the simulations gives overlap $\gamma^{-0.97}$ once the logarithm is divided out ($\gamma^{-1.21}$ as a plain power, against $\gamma^{-1.45}$ for the certified plan) and gap $\Delta^{-0.98}$ (certified $\Delta^{-1.38}$). With the gap held at 0.40 the cost grows as $N^{2.3}$; at fixed J_2, where g also closes, as $N^{3.1}$ (Fig. 14c). The difference is the $1/g$ of Eq. (C2).

Gap and overlap are not independent inputs. Tuning J_2 to hold the gap moves the chain toward the dimerized point $J_2 = 1/2$, where a $\chi = 2$ MPS is close to exact: at the gap-matched couplings the $\chi = 2$ overlap is 0.97–1.00 at every N, against 0.70–0.90 at $J_2 = 0.2411$. Up to $J_2 \approx 1/2$ a larger gap and a better trial come together, so for real trials the two effects compound; beyond it the gap keeps growing while the $\chi = 2$ overlap falls again (to 0.78 at $J_2 = 0.68$, $N = 14$). Holding one of them fixed through J_2 therefore always moves the other, which is why the scans here use synthetic trials and Sec. IV a sublattice rotation of a real trial circuit.

Appendix D: Reproducing the results

The code, raw results and paper sources are at `https://github.com/aaronseymour7/FilterPipeline`. After `pip install -r requirements.txt`, all commands below run from the repository root. `python sbpipe_study/run_study.py all` runs the DMRG front end (cached), the validation (A), scaling (B), sensitivity (C), ansatz (D) and oracle (E) studies and writes `results/*.json`; `python sbpipe_study/make_figures.py` writes every figure, table and the numerical macros used in this text. The stored results were produced with per-term CNOT counts and converted to the bond compile of Eq. (9) by `sbpipe_study/recompile_results.py`, which checks every stored count against the per-term formula before replacing it; a fresh run counts bonds directly. `python -m pytest sbpipe/tests` runs the component checks (Pauli algebra, Weyl bound, Trotter bounds against dense matrix exponentials, the Dirichlet identity (3), filter certificates against a dense grid, gate counts against an explicit merged sequence and against qiskit, and an end-to-end certificate check).

[1] I. Stetcu, A. Baroni, and J. Carlson, Projection algorithm for state preparation on quantum computers, Phys. Rev. C **108**, L031306 (2023), arXiv:2211.10545.

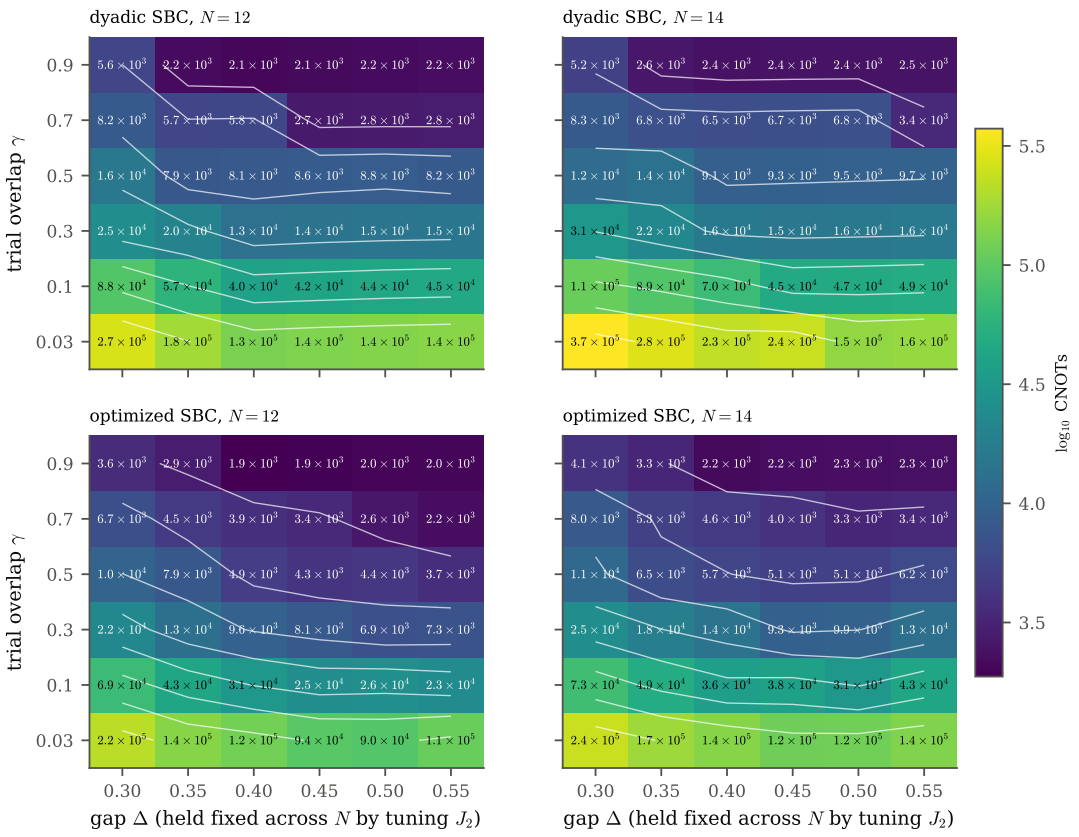


FIG. 12. Filter CNOTs per success over gap and overlap at $\epsilon = 10^{-2}$, measured, for the dyadic (top) and optimized (bottom) SBC filters at $N = 12$ and 14. White contours every factor 1.8. Source: `results/gapov.jsonl`, `results/gapov_grid.jsonl`.

[2] A. M. Childs, Y. Su, M. C. Tran, N. Wiebe, and S. Zhu, Theory of Trotter error with commutator scaling, Phys. Rev. X **11**, 011020 (2021).

[3] S. R. White, Density matrix formulation for quantum renormalization groups, Phys. Rev. Lett. **69**, 2863 (1992).

[4] U. Schollwöck, The density-matrix renormalization group in the age of matrix product states, Ann. Phys. **326**, 96 (2011).

[5] S.-J. Ran, Encoding of matrix product states into quantum circuits of one- and two-qubit gates, Phys. Rev. A **101**, 032310 (2020).

[6] Y. Ge, J. Tura, and J. I. Cirac, Faster ground state preparation and high-precision ground energy estimation with fewer qubits, J. Math. Phys. **60**, 022202 (2019).

[7] L. Lin and Y. Tong, Near-optimal ground state preparation, Quantum **4**, 372 (2020).

[8] A. Gilyén, Y. Su, G. H. Low, and N. Wiebe, Quantum singular value transformation and beyond, in *Proc. STOC 2019*, p. 193.

[9] Y. Dong, L. Lin, and Y. Tong, Ground-state preparation and energy estimation on early fault-tolerant quantum computers via quantum eigenvalue transformation of unitary matrices, PRX Quantum **3**, 040305 (2022).

[10] M. Bee-Lindgren, Z. Qian, M. DeCross, N. C. Brown, C. N. Gilbreth, J. Watkins, X. Zhang, and D. Lee, Controlled gate networks: theory and application to eigenvalue estimation, Eur. Phys. J. A (2025), doi:10.1140/epja/s10050-025-01750-y, arXiv:2208.13557.

[11] S. Fomichev, K. Hejazi, M. S. Zini, M. Kiser, J. Fraxanet, P. A. M. Casares, A. Delgado, J. Huh, A.-C. Voigt, J. E. Mueller, and J. M. Arrazola, Initial state preparation for quantum chemistry on quantum computers, PRX Quantum **5**, 040339 (2024).

[12] D. W. Berry, Y. Tong, T. Khattar, A. White, T. I. Kim, G. H. Low, S. Boixo, Z. Ding, L. Lin, S. Lee, G. K.-L. Chan, R. Babbush, and N. C. Rubin, Rapid initial-state preparation for the quantum simulation of strongly correlated molecules, PRX Quantum **6**, 020327 (2025).

[13] K. Choi, D. Lee, J. Bonitati, Z. Qian, and J. Watkins, Rodeo algorithm for quantum computing, Phys. Rev. Lett. **127**, 040505 (2021).

[14] N. J. Ross and P. Selinger, Optimal ancilla-free Clifford+T approximation of z-rotations, Quantum Inf. Comput. **16**, 901 (2016).

[15] V. V. Shende, S. S. Bullock, and I. L. Markov, Synthesis of quantum-logic circuits, IEEE Trans. Comput.-Aided Des. **25**, 1000 (2006).

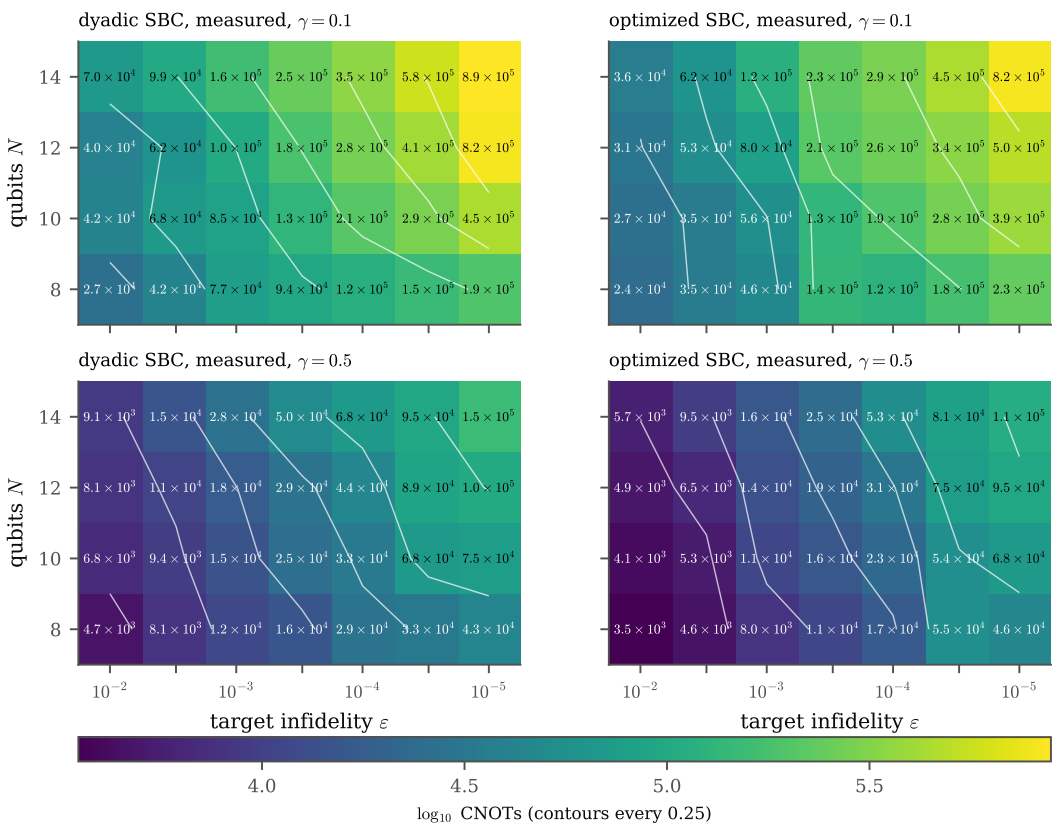


FIG. 13. Filter CNOTs per success over N and ϵ at a gap held at $g = 0.40$, measured, for the dyadic (left) and optimized (right) SBC filters at $\gamma = 0.1$ (top) and 0.5 (bottom). Source: `results/gapov*.jsonl`.

[16] K. Okamoto and K. Nomura, Fluid-dimer critical point in $S = 1/2$ antiferromagnetic Heisenberg chain with next nearest neighbor interactions, Phys. Lett. A **169**, 433 (1992).

[17] S. Eggert, Numerical evidence for multiplicative logarithmic corrections from marginal operators, Phys. Rev. B **54**, R9612 (1996).

[18] J. Gray, quimb: a python package for quantum information and many-body calculations, J. Open Source Softw. **3**, 819 (2018).

[19] B. Şahinoğlu and R. D. Somma, Hamiltonian simulation in the low-energy subspace, npj Quantum Inf. **7**, 119 (2021).

[20] Z. Qian, J. Watkins, G. Given, J. Bonitati, K. Choi, and D. Lee, Demonstration of the rodeo algorithm on a quantum computer, arXiv:2110.07747.

[21] E. Rule, I. Stetcu, and J. Carlson, Simplified projection on total spin zero for state preparation on quantum computers, arXiv:2410.02848 (2024).

[22] T. D. Cohen and H. Oh, Optimizing rodeo projection, arXiv:2305.19952 (2023).

[23] M. Patkowski, O. Ayyildiz, K. Hunt, N. Jansen, and D. Lee, Improved rodeo algorithm performance for spectral functions and state preparation, arXiv:2602.05978 (2026).

[24] C. F. Kane, N. Gomes, and M. Kreshchuk, Nearly-optimal state preparation for quantum simulations of lattice gauge theories, arXiv:2310.13757 (2023).

[25] T. R. Müller, M. Geiger, and C. B. Mendl, Ground-state preparation of the Fermi-Hubbard model on a quantum computer with 2D topology via quantum eigenvalue transformation of unitary matrices, arXiv:2411.18535 (2024).

[26] E. Karacan, Y. Chen, and C. B. Mendl, Enhancing scalability of quantum eigenvalue transformation of unitary matrices for ground state preparation through adaptive finer filtering, arXiv:2401.09091.

[27] G. Lee, M. Kang, J. Hong, S. Fomichev, and J. Huh, Filtered quantum phase estimation, arXiv:2510.04294 (2025).

[28] M. Watanabe, H. Nishi, T. Kosugi, S. Tsuneyuki, and Y. Matsushita, Overcoming the matrix-product-state encoding barrier via DMRG-guided probabilistic imaginary-time evolution, arXiv:2605.30141 (2026).

[29] N. Yoshioka, T. Okubo, Y. Suzuki, Y. Koizumi, and W. Mizukami, Hunting for quantum-classical crossover in condensed matter problems, npj Quantum Inf. (2024), arXiv:2210.14109.

[30] A. M. Childs, D. Maslov, Y. Nam, N. J. Ross, and Y. Su, Toward the first quantum simulation with quantum speedup, Proc. Natl. Acad. Sci. **115**, 9456 (2018).

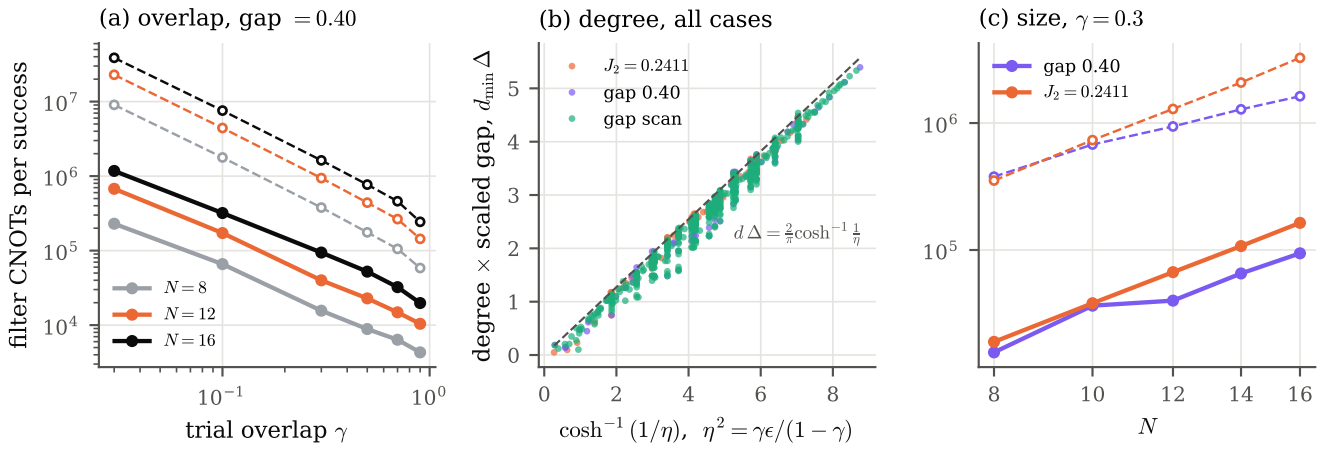


FIG. 14. What the filter actually needs, from exact spectra and state-vector simulation ($N = 8$–16, QETU, second order). (a) CNOTs per success versus overlap at a gap held at 0.40, $\epsilon = 10^{-3}$: simulated (filled) and certified (open, dashed). (b) Smallest degree times scaled gap for all 675 cases against $\cosh^{-1}(1/\eta)$; dashed: Eq. (C1). (c) Cost versus N at $\gamma = 0.3$, $\epsilon = 10^{-3}$, with the gap held at 0.40 or J_2 held at 0.2411. Source: `results/emp_trotter.json`, `sbpipe_study/run_empirical.py`.

[31] K. Mizuta and T. Kuwahara, Trotterization is substantially efficient for low-energy states, arXiv:2504.20746 (2025).

[32] F. Vatan and C. Williams, Optimal quantum circuits for general two-qubit gates, Phys. Rev. A **69**, 032315 (2004).

[33] G. Vidal and C. M. Dawson, Universal quantum circuit for two-qubit transformations with three controlled-NOT gates, Phys. Rev. A **69**, 010301 (2004).

[34] M. Amy, D. Maslov, and M. Mosca, Polynomial-time T-depth optimization of Clifford+T circuits via matroid partitioning, IEEE Trans. Comput.-Aided Des. **33**, 1476 (2014).

[35] G. Brassard, P. Høyer, M. Mosca, and A. Tapp, Quantum amplitude amplification and estimation, Contemp. Math. **305**, 53 (2002).

[36] T. J. Yoder, G. H. Low, and I. L. Chuang, Fixed-point quantum search with an optimal number of queries, Phys. Rev. Lett. **113**, 210501 (2014).

[37] R. Iten, R. Colbeck, I. Kukuljan, J. Home, and M. Christandl, Quantum circuits for isometries, Phys. Rev. A **93**, 032318 (2016).

[38] M. Plesch and Č. Brukner, Quantum-state preparation with universal gate decompositions, Phys. Rev. A **83**, 032302 (2011).

[39] J. A. Smolin and D. P. DiVincenzo, Five two-bit quantum gates are sufficient to implement the quantum Fredkin gate, Phys. Rev. A **53**, 2855 (1996).